

Lab: Computer Vision

ACTL3143 & ACTL5111 Deep Learning for Actuaries

Learning Goal

This lab uses one connected problem throughout: an insurer wants to classify vehicle inspection photos as either **damage** or **whole**.

The lab is grouped by topic. Each topic has numbered problems and a short concept explanation. Your job is to **work out the answers yourself** — by hand for the small calculations, and by completing the code stubs marked **# TODO** for the modelling parts.

! How to use this worksheet

- Each problem has a **Concept** box explaining the key idea. Read it if you are stuck, and revisit the Week 3 lecture notes for more detail.
- Code cells are intentionally left incomplete. Fill in the **# TODO** parts and run them in Jupyter.

Dataset used in the real-data section:

<https://www.kaggle.com/datasets/anujms/car-damage-detection>

Please download and unzip it into a **data/** folder beside this notebook, so the images live under **data/data1a/training/** and **data/data1a/validation/**.

```
import os
# Keras runs on the PyTorch backend in this course. This must be set before
# importing keras (otherwise Keras looks for TensorFlow and errors out).
os.environ["KERAS_BACKEND"] = "torch"

from pathlib import Path

import numpy as np
import matplotlib.pyplot as plt
```

```

from PIL import Image

import keras
from keras import layers

SEED = 3143
keras.utils.set_random_seed(SEED)

IMG_SIZE = (96, 128)      # height, width
BATCH_SIZE = 32
INPUT_SHAPE = IMG_SIZE + (3,)

# We force this order later so label 0 = whole and label 1 = damage.
DESIRED_CLASS_NAMES = ["01-whole", "00-damage"]

print("Keras:", keras.__version__, "| backend:", keras.config.backend())
print("Input shape:", INPUT_SHAPE)

```

```

Keras: 3.14.1 | backend: torch
Input shape: (96, 128, 3)

```

Topic 1: Images as Tensors

Problem 1.1: Pixel Encoding and Rescaling

An RGB pixel from one inspection photo has channel values written in hexadecimal as:

```
(0xB4, 0x98, 0x7E)
```

Convert it to decimal RGB values. Then divide each channel by 255, as a `Rescaling(1./255)` layer would do.

Concept

Each colour channel is stored as an integer from 0 to 255. Hexadecimal is just a compact way to write the same integer. CNNs usually train on rescaled channel values because inputs around $[0, 1]$ are numerically easier for gradient descent than raw values around $[0, 255]$.

```
hex_pixel = ("B4", "98", "7E")

# TODO: convert each hex string to a decimal integer (hint: int(channel, 16))
# TODO: build rgb_decimal as a numpy array, then rgb_rescaled = rgb_decimal / 255
# TODO: print the decimal RGB and the rescaled RGB (rounded to 3 dp)
```

Problem 1.2: Image and Batch Shapes

All images are resized to height 96, width 128, and 3 colour channels. A mini-batch contains 32 photos.

1. What is the shape of one image?
2. What is the shape of one mini-batch?
3. How many numeric channel values describe one image?
4. If you predict on one RGB image, what shape should be passed to Keras?

Concept

Keras Conv2D layers expect rank-4 input: (batch, height, width, channels). A single image still needs a batch axis, so a one-image prediction is a batch of length one.

```
# TODO: write expressions for each of the four answers using IMG_SIZE / INPUT_SHAPE / BATCH_SIZE
# one_image_shape = ...
# batch_shape = ...
# n_values = ... # hint: np.prod(...)
# one_prediction_shape = ... # hint: a leading axis of length 1
# print them all
```

Topic 2: A Local Convolution Filter

Problem 2.1: Compute One Filter Response

A small greyscale crop from a bumper photo is:

```
[[35, 42, 48],
 [40, 210, 52],
 [38, 45, 50]]
```

Use this local contrast filter with bias 0:

```
[[ -1, -1, -1],
 [ -1,  8, -1],
 [ -1, -1, -1]]
```

Compute the convolution output at the centre of the patch. Then apply ReLU.

💡 Concept

A convolution neuron takes a local patch, multiplies each pixel by the corresponding filter weight, adds the products, and adds a bias. The same filter is then reused across the image. This reuse is called weight sharing.

```
patch = np.array([
    [35, 42, 48],
    [40, 210, 52],
    [38, 45, 50],
])

contrast_filter = np.array([
    [-1, -1, -1],
    [-1,  8, -1],
    [-1, -1, -1],
])

# TODO: raw_response = elementwise product of patch and filter, summed (hint: np.sum(patch *
# TODO: relu_response = max(raw_response, 0)
# TODO: print both
```

Problem 2.2: Compare With a Flat Patch

Repeat the same calculation for a flat patch where all nine values are 60. What does this filter respond to? What kinds of damage cues would it pick up, and what could fool it?

💡 Concept

The same filter slides across the whole image (weight sharing). A filter that sums to zero responds to *local contrast* rather than overall brightness.

```
# TODO: build a 3x3 patch of all 60s (hint: np.full((3, 3), 60))
# TODO: compute and print its response with contrast_filter
# Then answer in a sentence: what does this filter respond to, and what could fool it?
```

Topic 3: CNN Shapes and Parameters

Problem 3.1: Trace a Scratch CNN

Consider this scratch CNN for 96x128 RGB photos:

```
Input((96, 128, 3))
Rescaling(1./255)
Conv2D(12, 3, padding="same", activation="relu")
MaxPooling2D()
Conv2D(24, 3, padding="same", activation="relu")
MaxPooling2D()
Conv2D(48, 3, padding="same", activation="relu")
MaxPooling2D()
Flatten()
Dense(32, activation="relu")
Dense(1)
```

Fill in the output shape and trainable parameter count after every trainable layer. Which single layer holds most of the parameters, and why?

Concept

For a convolution layer:

```
parameters = (kernel_height * kernel_width * input_channels + bias) * filters
```

For a dense layer:

```
parameters = (number_of_inputs + bias) * number_of_outputs
```

With `padding="same"` and stride 1, convolution keeps height and width unchanged. The default `MaxPooling2D()` halves each spatial dimension.

```
# The two formulas from the Concept box are provided as tools - use them to check your hand
def conv_params(kernel_size, in_channels, filters):
    return (kernel_size * kernel_size * in_channels + 1) * filters

def dense_params(n_in, n_out):
    return (n_in + 1) * n_out
```

```

# TODO: fill in (output_shape, params) for each row by hand, then verify with the helpers ab
layer_table = [
    ("Input", (96, 128, 3), 0),
    ("Rescaling", (96, 128, 3), 0),
    ("Conv2D(12, 3, same)", None, None),
    ("MaxPooling2D", None, 0),
    ("Conv2D(24, 3, same)", None, None),
    ("MaxPooling2D", None, 0),
    ("Conv2D(48, 3, same)", None, None),
    ("MaxPooling2D", None, 0),
    ("Flatten", None, 0),
    ("Dense(32)", (32,), None),
    ("Dense(1)", (1,), None),
]

# TODO: print the table and the total parameter count.

```

Problem 3.2: Replace Flatten With Global Average Pooling

Replace Flatten() by GlobalAveragePooling2D(), keeping Dense(32) -> Dense(1). Re-compute the trainable parameter count. Why does this reduce overfitting risk on a small image dataset?

Concept

GlobalAveragePooling2D() averages each feature map across height and width. A tensor of shape (12, 16, 48) becomes a vector of length 48, one summary per channel.

```

# TODO: keep the three conv layers, but feed Dense(32) from 48 channel means instead of 9,21
# gap_total = conv_params(...) + conv_params(...) + conv_params(...) + dense_params(48, 32)
# TODO: print gap_total and compare it with the Flatten total from Problem 3.1.

```

Problem 3.3: Build the Scratch Model in Keras

Implement the architecture above as a Keras functional model, so you can check your hand calculations against model.summary(). The same builder is reused in later topics, so support optional augmentation, global average pooling, and dropout via arguments.

💡 Concept

The final output has one unit and **no sigmoid**, because the model will be trained with `BinaryCrossentropy(from_logits=True)`.

```
def build_scratch_cnn(input_shape=INPUT_SHAPE, use_gap=False, use_augmentation=False, dropout=0.5):
    inputs = keras.Input(shape=input_shape)
    x = inputs

    # TODO (optional, used later): if use_augmentation, add a small augmentation block
    # (RandomFlip("horizontal"), RandomRotation(0.04), RandomZoom(0.08)).

    # TODO: Rescaling(1/255)
    # TODO: Conv2D(12, 3, padding="same", activation="relu") then MaxPooling2D()
    # TODO: Conv2D(24, 3, padding="same", activation="relu") then MaxPooling2D()
    # TODO: Conv2D(48, 3, padding="same", activation="relu") then MaxPooling2D()
    # TODO: if use_gap -> GlobalAveragePooling2D(), else -> Flatten()
    # TODO: if dropout > 0 -> Dropout(dropout)
    # TODO: Dense(32, activation="relu")
    # TODO: outputs = Dense(1) (a single logit, no activation)

    outputs = x # replace with your Dense(1) output
    return keras.Model(inputs, outputs, name="scratch_vehicle_cnn")

# scratch_shape_model = build_scratch_cnn(use_gap=False)
# scratch_shape_model.summary() # compare the Param # column against Problems 3.1 and 3.2
```

Topic 4: Binary Logits and Evaluation

Problem 4.1: Convert Logits to Probabilities

The final `Dense(1)` returns one logit z . A positive logit favours **damage**; a negative logit favours **whole**.

The probability of damage is:

$$\text{sigmoid}(z) = 1 / (1 + \exp(-z))$$

Convert $z = 1.2$ and $z = -0.8$ to probabilities. Which class does each predict? If accuracy is computed directly on logits, what is the decision threshold — 0.0 or 0.5?

💡 Concept

`sigmoid(0) = 0.5`. So a probability threshold of 0.5 corresponds to a logit threshold of 0.0.

```
logits = np.array([1.2, -0.8])
```

```
# TODO: prob_damage = 1 / (1 + np.exp(-logits))
```

```
# TODO: for each logit, decide the predicted class (damage if z >= 0 else whole) and print z
```

Problem 4.2: Compile a Binary-Logit Model

Write a helper that compiles any of these single-logit models with the right loss and metrics: `BinaryCrossentropy(from_logits=True)`, an accuracy that thresholds the logit at 0.0, and `AUC(from_logits=True)`.

💡 Concept

`BinaryCrossentropy(from_logits=True)` tells Keras to apply the sigmoid inside the loss in a numerically stable way. For accuracy computed directly on logits, the threshold is 0.0.

Some Keras versions require `BinaryAccuracy.threshold` to be strictly between 0 and 1, so it is cleaner to define a small custom logit-threshold metric.

```
def binary_logit_accuracy(y_true, y_pred):
    # TODO: cast y_true to float32
    # TODO: predicted class = 1.0 where y_pred (a logit) >= 0.0, else 0.0
    # TODO: return 1.0 where predicted class == y_true, else 0.0 (use keras.ops)
    raise NotImplementedError

def compile_binary_logit_model(model, lr=1e-3):
    # TODO: model.compile(
    #     optimizer = Adam(learning_rate=lr),
    #     loss      = BinaryCrossentropy(from_logits=True),
    #     metrics   = [binary_logit_accuracy, AUC(from_logits=True, name="auc")],
    # )
    return model
```

```
# demo_model = build_scratch_cnn(use_gap=True)
# compile_binary_logit_model(demo_model, lr=1e-3)
```

Topic 5: Real Kaggle Data

The remaining topics give you a working data pipeline so you can focus on understanding and running it. Read each cell, run it, and answer the questions in the markdown.

Problem 5.1: Download or Locate the Dataset

Use KaggleHub to download the dataset if it is not already present.

Concept

Kaggle datasets usually do not have a stable unauthenticated direct ZIP URL. KaggleHub is the clean Python interface:

```
import kagglehub
path = kagglehub.dataset_download("anujms/car-damage-detection")
```

Here we can add `output_dir` so the files stay inside the local folder.

```
DATASET_HANDLE = "anujms/car-damage-detection"

# The dataset has been downloaded and unzipped into ./data/data1a.
DATA_DIR = Path("data") / "data1a"
TRAIN_DIR = DATA_DIR / "training"
VAL_DIR = DATA_DIR / "validation"

# If you do not have the data yet, download it from Kaggle with KaggleHub:
#     import kagglehub
#     kagglehub.dataset_download(DATASET_HANDLE)
# then move the data1a folder to ./data/data1a.

print("DATA_DIR:", DATA_DIR)
```

Problem 5.2: Count the Images

Run the cell and write down the class balance. Is the dataset balanced? How would your choice of evaluation metric change if it were heavily imbalanced?

💡 Concept

Before modelling, check the dataset layout and class balance. A balanced dataset makes accuracy easier to interpret; an imbalanced dataset requires more care with AUC, recall, precision, and class weights.

```
IMAGE_EXTENSIONS = {".jpg", ".jpeg", ".png", ".bmp", ".gif"}

def count_images(folder):
    return sum(1 for p in folder.rglob("*") if p.suffix.lower() in IMAGE_EXTENSIONS)

for split_dir in [TRAIN_DIR, VAL_DIR]:
    print(f"\n{split_dir.name}")
    for class_dir in sorted(p for p in split_dir.iterdir() if p.is_dir()):
        print(f"  {class_dir.name:10s} {count_images(class_dir):5d} images")
```

Problem 5.3: Load and Inspect Batches

Load the images straight into NumPy arrays with a small helper. We pass the class order ["01-whole", "00-damage"] so that the label is the *position* in that list:

```
whole -> 0
damage -> 1
```

Question: what label would `damage` get if you instead ordered the class folders *alphabetically* (00-damage, 01-whole)? Why might that be confusing given the usual positive-class convention?

💡 Concept

We load each image with Pillow, resize it to `IMG_SIZE`, and stack the results into an array of shape `(n, height, width, 3)`. The label is the index of the class in the list we pass in, so choosing the order explicitly keeps `damage` as the positive class (label 1). We keep the raw 0-255 pixel values because the model's `Rescaling(1./255)` layer does the scaling.

```
def load_images_from_directory(directory, class_names, img_size=IMG_SIZE):
    """Load every image under directory/<class>/ into NumPy arrays.
```

Returns X of shape (n, height, width, 3) as uint8, and y of shape (n, 1) as float32, where each label is the index of its class in ``class\_names`` (so whole -> 0 and damage -> 1).

```
"""
directory = Path(directory)
height, width = img_size
paths, labels = [], []
for index, class_name in enumerate(class_names):
    for path in sorted((directory / class_name).iterdir()):
        if path.suffix.lower() in IMAGE_EXTENSIONS:
            paths.append(path)
            labels.append(index)

X = np.empty((len(paths), height, width, 3), dtype=np.uint8)
for i, path in enumerate(paths):
    img = Image.open(path).convert("RGB").resize((width, height), Image.BILINEAR)
    X[i] = np.asarray(img, dtype=np.uint8)
y = np.array(labels, dtype=np.float32).reshape(-1, 1)
return X, y
```

```
class_names = DESIRED_CLASS_NAMES
X_train, y_train = load_images_from_directory(TRAIN_DIR, class_names)
X_val, y_val = load_images_from_directory(VAL_DIR, class_names)

# Shuffle once so batches (and the small samples below) mix the two classes.
rng = np.random.default_rng(SEED)
train_perm = rng.permutation(len(X_train))
X_train, y_train = X_train[train_perm], y_train[train_perm]
val_perm = rng.permutation(len(X_val))
X_val, y_val = X_val[val_perm], y_val[val_perm]

print("class_names:", class_names)
print("X_train:", X_train.shape, "y_train:", y_train.shape)
print("X_val:", X_val.shape, "y_val:", y_val.shape)
```

```
plt.figure(figsize=(9, 5))
for i in range(8):
    ax = plt.subplot(2, 4, i + 1)
    ax.imshow(X_train[i])
    label_index = int(y_train[i, 0])
```

```
ax.set_title(class_names[label_index])
ax.axis("off")
plt.tight_layout()
```

Topic 6: Scratch CNN on Real Images

Problem 6.1: Train a Small Sanity-Check Model

Train the global-average-pooling version on a few batches only. This is a **pipeline check**, not a final performance estimate. (Requires your completed `build_scratch_cnn` and `compile_binary_logit_model` from Topics 3–4.)

Concept

A short sanity run answers: “Does the data pipeline, model, loss, and metric setup work?”
It does not answer: “What is the best achievable validation AUC?”

```
# Use a few batches' worth of images as a quick pipeline check.
X_train_small, y_train_small = X_train[:6 * BATCH_SIZE], y_train[:6 * BATCH_SIZE]
X_val_small, y_val_small = X_val[:3 * BATCH_SIZE], y_val[:3 * BATCH_SIZE]

sanity_model = build_scratch_cnn(use_gap=True, use_augmentation=True, dropout=0.25)
compile_binary_logit_model(sanity_model, lr=5e-4)

sanity_history = sanity_model.fit(
    X_train_small,
    y_train_small,
    validation_data=(X_val_small, y_val_small),
    batch_size=BATCH_SIZE,
    epochs=2,
    verbose=2,
)
```

Problem 6.2: Confusion Matrix From Logits

Collect logits over the validation sample, threshold at 0.0, and build a confusion matrix with precision and recall for the **damage** class. Why do we use logits for the loss but convert to probabilities for interpretation?

💡 Concept

Use logits for the loss, but convert to probabilities for interpretation. A logit threshold of 0.0 is equivalent to a probability threshold of 0.5.

```
def collect_logits_and_labels(model, X, y):
    logits = model.predict(X, verbose=0).reshape(-1)
    labels = y.reshape(-1).astype(int)
    return logits, labels

def binary_report_from_logits(logits, labels):
    probs = 1 / (1 + np.exp(-logits))
    preds = (logits >= 0).astype(int)
    cm = np.zeros((2, 2), dtype=int)
    for y_true, y_pred in zip(labels, preds):
        cm[y_true, y_pred] += 1

    accuracy_value = np.mean(preds == labels)
    precision_damage = cm[1, 1] / max(cm[0, 1] + cm[1, 1], 1)
    recall_damage = cm[1, 1] / max(cm[1, 0] + cm[1, 1], 1)

    print("confusion matrix rows=true [whole, damage], columns=predicted [whole, damage]")
    print(cm)
    print(f"accuracy: {accuracy_value:.3f}")
    print(f"damage precision: {precision_damage:.3f}")
    print(f"damage recall: {recall_damage:.3f}")
    print("first five probabilities:", np.round(probs[:5], 3).tolist())
    return probs, preds, cm

sanity_logits, sanity_labels = collect_logits_and_labels(sanity_model, X_val_small, y_val_small)
sanity_probs, sanity_preds, sanity_cm = binary_report_from_logits(sanity_logits, sanity_labels)
```

Problem 6.3: What Mistakes Should We Expect?

Before looking at the wrong-prediction plot below, **predict** what kinds of images this model is likely to misclassify, and write down at least four. Then run the cell and compare your list against what you actually see.

```

logits = sanity_model.predict(X_val_small, verbose=0).reshape(-1)
preds = (logits >= 0).astype(int)
labels = y_val_small.reshape(-1).astype(int)

shown = 0
plt.figure(figsize=(9, 5))
for image, y_true, y_pred, z in zip(X_val_small, labels, preds, logits):
    if y_true != y_pred and shown < 8:
        ax = plt.subplot(2, 4, shown + 1)
        ax.imshow(image)
        ax.set_title(f"true={class_names[y_true]}\npred={class_names[y_pred]}\nz={z:.2f}")
        ax.axis("off")
        shown += 1
    if shown >= 8:
        break
plt.tight_layout()
if shown == 0:
    print("No wrong predictions found in this small validation sample.")

```

Topic 7: Stretch Refinements and Transfer Learning

Problem 7.1: Which Refinement Changes the Parameter Count Most?

Of the common refinements — data augmentation, dropout, and replacing `Flatten()` with `GlobalAveragePooling2D()` — which one changes the **trainable parameter count** the most, and which mainly help generalisation without adding parameters? Justify your answer using the counts from Topic 3, then build the refined model and confirm with `summary()`.

Concept

Augmentation and dropout act on activations, not weights, so they add (almost) no trainable parameters. Structural changes to the head are what move the parameter count.

```

refined_model = build_scratch_cnn(use_gap=True, use_augmentation=True, dropout=0.25)
compile_binary_logit_model(refined_model, lr=5e-4)
refined_model.summary()

```

Problem 7.2: Full Training Code

The following is the full training version. Run it when you have time and the dataset available. Compare its validation AUC/loss with the sanity run.

```
callbacks = [  
    keras.callbacks.EarlyStopping(  
        monitor="val_loss",  
        patience=3,  
        restore_best_weights=True,  
    )  
]  
  
history = refined_model.fit(  
    X_train,  
    y_train,  
    validation_data=(X_val, y_val),  
    batch_size=BATCH_SIZE,  
    epochs=12,  
    callbacks=callbacks,  
)
```

Problem 7.3: Optional Transfer Learning With MobileNetV2

Build a transfer-learning model on a frozen MobileNetV2 base with a small binary head.

Question: why is transfer learning often more reliable than training from scratch on a small specialised dataset, and why should fine-tuning use a very small learning rate (e.g. $1e-5$) on only the last block?

Concept

Transfer learning reuses a CNN trained on a large general image dataset. The early and middle layers already know visual features such as edges, textures, corners, and object parts. On a small specialised dataset, training only a small head is often more reliable than training every convolutional weight from scratch.

```
def build_mobilenet_transfer_model(input_shape=INPUT_SHAPE, dropout=0.25):  
    base = keras.applications.MobileNetV2(  
        input_shape=input_shape,  
        include_top=False,  
        weights="imagenet",
```

```

)
base.trainable = False

inputs = keras.Input(shape=input_shape)
x = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.04),
        layers.RandomZoom(0.08),
    ],
    name="transfer_augmentation",
)(inputs)
x = keras.applications.mobilenet_v2.preprocess_input(x)
x = base(x, training=False)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(dropout)(x)
outputs = layers.Dense(1, name="damage_logit")(x)
model = keras.Model(inputs, outputs, name="mobilenetv2_vehicle_damage")
return model, base

transfer_model, transfer_base = build_mobilenet_transfer_model()
compile_binary_logit_model(transfer_model, lr=1e-3)

transfer_history = transfer_model.fit(
    X_train,
    y_train,
    validation_data=(X_val, y_val),
    batch_size=BATCH_SIZE,
    epochs=8,
    callbacks=callbacks,
)

# Optional fine-tuning.
transfer_base.trainable = True
for layer in transfer_base.layers[:-20]:
    layer.trainable = False

compile_binary_logit_model(transfer_model, lr=1e-5)

fine_tune_history = transfer_model.fit(
    X_train,

```

```
y_train,  
validation_data=(X_val, y_val),  
batch_size=BATCH_SIZE,  
epochs=5,  
callbacks=callbacks,  
)
```

Closing Checklist

By the end of this lab, you should be able to:

1. Read RGB image tensors and Keras batch shapes.
2. Compute one convolution response by hand.
3. Trace CNN output shapes and parameter counts.
4. Explain why `Flatten()` can create many parameters.
5. Interpret binary logits, sigmoid probabilities, and thresholding.
6. Load a real image-folder dataset with explicit class ordering.
7. Train and evaluate a scratch CNN pipeline.
8. Explain why data augmentation, dropout, global average pooling, and transfer learning can matter.