

# Classification & Optimisation

ACTL3143 & ACTL5111 Deep Learning for Actuaries

Patrick Laub

# Overview

In these slides, we'll start by giving some demonstrations of training classification models that: 1) predict a binary outcome, then 2) predict a categorical outcome with  $> 2$  options or levels.

Next, we'll step into the maths of how these classification models make predictions, then go look at the high-level ideas of how to “train” them, then finally look at the maths of this training process.

# Imports needed for these demos

```
1 import random
2 from pathlib import Path
3
4 import matplotlib.pyplot as plt
5 import numpy as np
6 import pandas as pd
7 import seaborn as sns
8
9 from keras.models import Sequential
10 from keras.layers import Dense, Input
11 from keras.callbacks import EarlyStopping
12
13 from sklearn.datasets import load_iris
14 from sklearn.model_selection import train_test_split
15 from sklearn.preprocessing import OneHotEncoder, StandardScaler
16 from sklearn.compose import make_column_transformer
17 from sklearn.impute import SimpleImputer
18 from sklearn.metrics import confusion_matrix, RocCurveDisplay, PrecisionRecallDisplay
19 from sklearn import set_config
20
21 set_config(transform_output="pandas")
```

# Lecture Outline

- **Binary Classification**
- Multiclass Classification
- Dense Layers in Matrices
- Optimisation
- Loss and Derivatives

# Stroke Prediction Data description

1. **id**: unique identifier
2. **gender**: “Male”, “Female” or “Other”
3. **age**: age of the patient
4. **hypertension**: 0 or 1 if the patient has hypertension
5. **heart\_disease**: 0 or 1 if the patient has any heart disease
6. **ever\_married**: “No” or “Yes”
7. **work\_type**: “children”, “Govt\_jov”, “Never\_worked”, “Private” or “Self-employed”
8. **Residence\_type**: “Rural” or “Urban”
9. **avg\_glucose\_level**: average glucose level in blood
10. **bmi**: body mass index
11. **smoking\_status**: “formerly smoked”, “never smoked”, “smokes” or “Unknown”
12. **stroke**: 0 or 1 if the patient had a stroke

Source: Kaggle, **Stroke Prediction Dataset**.

# Load up the (pre-)preprocessed data

```

1 PROCESSED_DATA_DIR = Path("stroke/processed")
2
3 X_train = pd.read_csv(PROCESSED_DATA_DIR / "x_train.csv")
4 X_val= pd.read_csv(PROCESSED_DATA_DIR / "x_val.csv")
5 X_test = pd.read_csv(PROCESSED_DATA_DIR / "x_test.csv")
6 y_train = pd.read_csv(PROCESSED_DATA_DIR / "y_train.csv")
7 y_val = pd.read_csv(PROCESSED_DATA_DIR / "y_val.csv")
8 y_test = pd.read_csv(PROCESSED_DATA_DIR / "y_test.csv")
9
10 X_train

```

|             | gender_Female | gender_Male | ever_married_No | ever_married_Yes | Residence_type_Rural | Residence_type_Urban |
|-------------|---------------|-------------|-----------------|------------------|----------------------|----------------------|
| <b>0</b>    | 0.0           | 1.0         | 0.0             | 1.0              | 1.0                  | 0.0                  |
| <b>1</b>    | 0.0           | 1.0         | 1.0             | 0.0              | 1.0                  | 0.0                  |
| <b>2</b>    | 0.0           | 1.0         | 1.0             | 0.0              | 1.0                  | 0.0                  |
| <b>...</b>  | ...           | ...         | ...             | ...              | ...                  | ...                  |
| <b>3063</b> | 1.0           | 0.0         | 0.0             | 1.0              | 1.0                  | 0.0                  |
| <b>3064</b> | 1.0           | 0.0         | 0.0             | 1.0              | 1.0                  | 0.0                  |
| <b>3065</b> | 0.0           | 1.0         | 1.0             | 0.0              | 0.0                  | 1.0                  |

3066 rows × 20 columns

# Target variable

```
1 y_train
```

|             | stroke     |
|-------------|------------|
| <b>0</b>    | 0          |
| <b>1</b>    | 0          |
| <b>2</b>    | 0          |
| <b>...</b>  | <b>...</b> |
| <b>3063</b> | 0          |
| <b>3064</b> | 0          |
| <b>3065</b> | 0          |

3066 rows  $\times$  1 columns

```
1 classes, counts = np.unique(y_train.values.ravel(), return_counts=True)
2 print("Classes:", classes)
3 print("Counts:", counts)
```

Classes: [0 1]  
Counts: [2909 157]

# Setup a binary classification model

```

1 def create_model(seed=42):
2     random.seed(seed)
3     model = Sequential()
4     model.add(Input(X_train.shape[1:]))
5     model.add(Dense(32, "leaky_relu"))
6     model.add(Dense(16, "leaky_relu"))
7     model.add(Dense(1, "sigmoid"))
8     return model

```

```

1 model = create_model()
2 model.summary()

```

Model: "sequential"

| Layer (type)    | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense)   | (None, 32)   | 672     |
| dense_1 (Dense) | (None, 16)   | 528     |
| dense_2 (Dense) | (None, 1)    | 17      |

Total params: 1,217 (4.75 KB)

Trainable params: 1,217 (4.75 KB)

Non-trainable params: 0 (0.00 B)



# Fit the model

```
1 model = create_model()  
2 model.compile("adam", "binary_crossentropy")  
3 model.fit(X_train, y_train, epochs=5, verbose=2)
```

Epoch 1/5

96/96 - 0s - 2ms/step - loss: 0.2734

Epoch 2/5

96/96 - 0s - 2ms/step - loss: 0.1753

Epoch 3/5

96/96 - 0s - 2ms/step - loss: 0.1665

Epoch 4/5

96/96 - 0s - 2ms/step - loss: 0.1619

Epoch 5/5

96/96 - 0s - 2ms/step - loss: 0.1595

<keras.src.callbacks.history.History at 0x127102900>

# Track accuracy as the model trains

```
1 model = create_model()  
2 model.compile("adam", "binary_crossentropy", metrics=["accuracy"])  
3 model.fit(X_train, y_train, epochs=5, verbose=2)
```

Epoch 1/5

96/96 - 0s - 2ms/step - accuracy: 0.9204 - loss: 0.2711

Epoch 2/5

96/96 - 0s - 2ms/step - accuracy: 0.9488 - loss: 0.1766

Epoch 3/5

96/96 - 0s - 2ms/step - accuracy: 0.9488 - loss: 0.1667

Epoch 4/5

96/96 - 0s - 2ms/step - accuracy: 0.9488 - loss: 0.1623

Epoch 5/5

96/96 - 0s - 2ms/step - accuracy: 0.9488 - loss: 0.1595

<keras.src.callbacks.history.History at 0x126e8afd0>

# Run a long fit

```
1 model = create_model()  
2 model.compile("adam", "binary_crossentropy", metrics=["accuracy"])  
3 %time hist = model.fit(X_train, y_train, epochs=500, validation_data=(X_val, y_val), verb
```

CPU times: user 2min 3s, sys: 9.19 s, total: 2min 12s

Wall time: 2min 6s

# Add early stopping



```
1 model = create_model()
2 model.compile("adam", "binary_crossentropy", metrics=["accuracy"])
3 es = EarlyStopping(restore_best_weights=True, patience=50, monitor="val_accuracy")
4 %time hist_es = model.fit(X_train, y_train, epochs=500, validation_data=(X_val, y_val), c
5 print(f"Stopped after {len(hist_es.history['loss'])} epochs.")
```

CPU times: user 12.7 s, sys: 940 ms, total: 13.6 s

Wall time: 13 s

Stopped after 51 epochs.

# Fitting metrics



# Add metrics, compile, and fit

```
1 model = create_model()
2
3 pr_auc = keras.metrics.AUC(curve="PR", name="pr_auc")
4 model.compile(optimizer="adam", loss="binary_crossentropy",
5               metrics=[pr_auc, "accuracy", "auc"])
6
7 es = EarlyStopping(patience=50, restore_best_weights=True,
8                   monitor="val_pr_auc", verbose=1)
9 model.fit(X_train, y_train, callbacks=[es], epochs=1_000, verbose=0,
10         validation_data=(X_val, y_val));
```

Epoch 81: early stopping


Restoring model weights from the end of the best epoch: 31.

```
1 model.evaluate(X_val, y_val, verbose=0)
```

```
[0.14898666739463806,
0.12857568264007568,
0.9569471478462219,
0.8119411468505859]
```

# Why use cross-entropy loss?

```
1 p = np.linspace(0, 1, 100)
2 plt.plot(p, (1 - p) ** 2)
3 plt.plot(p, -np.log(p))
4 plt.legend(["MSE", "Cross-entropy"]);
```



# Overweight the minority class

```

1 model = create_model()
2
3 pr_auc = keras.metrics.AUC(curve="PR", name="pr_auc")
4 model.compile(optimizer="adam", loss="binary_crossentropy",
5               metrics=[pr_auc, "accuracy", "auc"])
6
7 es = EarlyStopping(patience=50, restore_best_weights=True,
8                   monitor="val_pr_auc", verbose=1)
9 model.fit(X_train, y_train.to_numpy(), callbacks=[es], epochs=1_000, verbose=0,
10         validation_data=(X_val, y_val), class_weight={0: 1, 1: 10});

```

Epoch 64: early stopping

Restoring model weights from the end of the best epoch: 14.

```
1 model.evaluate(X_val, y_val, verbose=0)
```

```

[0.3523019552230835,
 0.13380154967308044,
 0.7896282076835632,
 0.8259596824645996]

```

```
1 model.evaluate(X_test, y_test, verbose=0)
```

```

[0.36996063590049744,
 0.15842117369174957,
 0.7954990267753601,
 0.8060390949249268]


```



# Classification Metrics

```
1 y_pred = model.predict(X_test, verbose=0)
```


```
1 RocCurveDisplay.from_predictions(y_test, y_pred, name="");
```



```
1 y_pred_stroke = y_pred > 0.5
2 confusion_matrix(y_test, y_pred_stroke)
```

```
array([[778, 194],
       [ 15, 35]])
```

```
1 PrecisionRecallDisplay.from_predictions(y_test, y_pred, name="");
```



```
1 y_pred_stroke = y_pred > 0.3
2 confusion_matrix(y_test, y_pred_stroke)
```

```
array([[647, 325],
       [ 7, 43]])
```

# Lecture Outline

- Binary Classification
- **Multiclass Classification**
- Dense Layers in Matrices
- Optimisation
- Loss and Derivatives

# Iris dataset

```
1 iris = load_iris()
2 names = ["SepalLength", "SepalWidth", "PetalLength", "PetalWidth"]
3 features = pd.DataFrame(iris.data, columns=names)
4 features
```

|            | SepalLength | SepalWidth | PetalLength | PetalWidth |
|------------|-------------|------------|-------------|------------|
| <b>0</b>   | 5.1         | 3.5        | 1.4         | 0.2        |
| <b>1</b>   | 4.9         | 3.0        | 1.4         | 0.2        |
| <b>...</b> | ...         | ...        | ...         | ...        |
| <b>148</b> | 6.2         | 3.4        | 5.4         | 2.3        |
| <b>149</b> | 5.9         | 3.0        | 5.1         | 1.8        |

150 rows × 4 columns

# Target variable

```
1 iris.target_names
```

```
array(['setosa', 'versicolor', 'virginica'],
      dtype='<U10')
```

```
1 iris.target[:8]
```

```
array([0, 0, 0, 0, 0, 0, 0, 0])
```

```
1 target = iris.target
2 target = target.reshape(-1, 1)
3 target[:8]
```

```
array([[0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0]])
```

```
1 classes, counts = np.unique(
2     target,
3     return_counts=True
4 )
5 print(classes)
6 print(counts)
```

```
[0 1 2]
[50 50 50]
```

```
1 iris.target_names[
2     target[[0, 30, 60]]
3 ]
```

```
array(['setosa',
       'setosa',
       'versicolor'], dtype='<U10')
```

# Split the data into train and test

```
1 X_train, X_test, y_train, y_test = train_test_split(features, target, random_state=24)
2 X_train
```


|            | SepalLength | SepalWidth | PetalLength | PetalWidth |
|------------|-------------|------------|-------------|------------|
| <b>53</b>  | 5.5         | 2.3        | 4.0         | 1.3        |
| <b>58</b>  | 6.6         | 2.9        | 4.6         | 1.3        |
| <b>95</b>  | 5.7         | 3.0        | 4.2         | 1.2        |
| <b>...</b> | <b>...</b>  | <b>...</b> | <b>...</b>  | <b>...</b> |
| <b>145</b> | 6.7         | 3.0        | 5.2         | 2.3        |
| <b>87</b>  | 6.3         | 2.3        | 4.4         | 1.3        |
| <b>131</b> | 7.9         | 3.8        | 6.4         | 2.0        |

112 rows  $\times$  4 columns

```
1 X_test.shape, y_test.shape
```

((38, 4), (38, 1))

# A basic classifier network



A basic network for classifying into three categories.

Source: Marcus Lautier (2022).

# Create a classifier model

```
1 NUM_FEATURES = len(features.columns)
2 NUM_CATS = len(np.unique(target))
3
4 print("Number of features:", NUM_FEATURES)
5 print("Number of categories:", NUM_CATS)
```

Number of features: 4

Number of categories: 3

Make a function to return a Keras model:

```
1 def build_model(seed=42):
2     random.seed(seed)
3     return Sequential([
4         Dense(30, activation="relu"),
5         Dense(NUM_CATS, activation="softmax")
6     ])
```

# Fit the model

```
1 model = build_model()  
2 model.compile("adam", "sparse_categorical_crossentropy")  
3  
4 model.fit(X_train, y_train, epochs=5, verbose=2);
```

Epoch 1/5

4/4 - 0s - 2ms/step - loss: 1.3920

Epoch 2/5

4/4 - 0s - 2ms/step - loss: 1.2912

Epoch 3/5

4/4 - 0s - 2ms/step - loss: 1.2196

Epoch 4/5

4/4 - 0s - 2ms/step - loss: 1.1576

Epoch 5/5

4/4 - 0s - 2ms/step - loss: 1.1084



# Track accuracy as the model trains

```
1 model = build_model()  
2 model.compile("adam", "sparse_categorical_crossentropy", metrics=["accuracy"])  
3 model.fit(X_train, y_train, epochs=5, verbose=2);
```

Epoch 1/5

4/4 - 0s - 2ms/step - accuracy: 0.2857 - loss: 1.3930

Epoch 2/5

4/4 - 0s - 2ms/step - accuracy: 0.2857 - loss: 1.2970

Epoch 3/5

4/4 - 0s - 2ms/step - accuracy: 0.2857 - loss: 1.2203

Epoch 4/5

4/4 - 0s - 2ms/step - accuracy: 0.2946 - loss: 1.1596

Epoch 5/5

4/4 - 0s - 2ms/step - accuracy: 0.3393 - loss: 1.1067

# Run a long fit

```
1 model = build_model()  
2 model.compile("adam", "sparse_categorical_crossentropy", \  
3     metrics=["accuracy"])  
4 %time hist = model.fit(X_train, y_train, epochs=500, \  
5     validation_split=0.25, verbose=False)
```

CPU times: user 3.84 s, sys: 466 ms, total: 4.31 s

Wall time: 4 s

Evaluation now returns both *loss* and *accuracy*.

```
1 model.evaluate(X_test, y_test, verbose=False)
```

[0.08639740198850632, 0.9736841917037964]

# Add early stopping

```

1 model_es = build_model()
2 model_es.compile("adam", "sparse_categorical_crossentropy", \
3                 metrics=["accuracy"])
4
5 es = EarlyStopping(restore_best_weights=True, patience=50,
6                   monitor="val_accuracy")
7 %time hist_es = model_es.fit(X_train, y_train, epochs=500, \
8                             validation_split=0.25, callbacks=[es], verbose=False);
9
10 print(f"Stopped after {len(hist_es.history['loss'])} epochs.")

```

CPU times: user 533 ms, sys: 65.6 ms, total: 599 ms

Wall time: 555 ms

Stopped after 70 epochs.

## Evaluation on test set:



```

1 model_es.evaluate(X_test, y_test, verbose=False)

```

[0.8077937960624695, 0.9210526347160339]

# Fitting metrics



# What is the softmax activation?

It creates a “probability” vector:  $\text{Softmax}(\mathbf{x})_i = \frac{e^{x_i}}{\sum_j e^{x_j}}$ .

In NumPy:

```
1 out = np.array([5, -1, 6])  
2 (np.exp(out) / np.exp(out).sum()).round(3)
```

```
array([0.27, 0.  , 0.73])
```

In Keras:

```
1 out = keras.ops.convert_to_tensor([[5.0, -1.0, 6.0]])  
2 keras.ops.round(keras.ops.softmax(out), 3)
```

```
tensor([[0.2690, 0.0010, 0.7310]])
```

# Prediction using classifiers

```
1 y_test[:4]
```

```
array([[2],
       [2],
       [1],
       [1]])
```

```
1 y_pred = model.predict(X_test.head(4), verbose=0)
2 y_pred
```

```
array([[2.02e-06, 7.64e-02, 9.24e-01],
       [1.86e-07, 1.62e-03, 9.98e-01],
       [1.44e-02, 9.76e-01, 1.00e-02],
       [2.80e-03, 8.50e-01, 1.48e-01]], dtype=float32)
```

```
1 # Add 'keepdims=True' to get a column vector.
2 np.argmax(y_pred, axis=1)
```

```
array([2, 2, 1, 1])
```

```
1 iris.target_names[np.argmax(y_pred, axis=1)]
```

```
array(['virginica', 'virginica', 'versicolor', 'versicolor'], dtype='<U10')
```

# Summary: Classification models in Keras

If the number of classes is  $c$ , then:

| Target                     | Output Layer                                        | Loss Function             |
|----------------------------|-----------------------------------------------------|---------------------------|
| Binary<br>( $c = 2$ )      | 1 neuron with <code>sigmoid</code><br>activation    | Binary Cross-Entropy      |
| Multi-class<br>( $c > 2$ ) | $c$ neurons with <code>softmax</code><br>activation | Categorical Cross-Entropy |

# Summary: Optionally output logits

If the number of classes is  $c$ , then:

| Target                     | Output Layer                                       | Loss Function                                                  |
|----------------------------|----------------------------------------------------|----------------------------------------------------------------|
| Binary<br>( $c = 2$ )      | 1 neuron with <code>linear</code><br>activation    | Binary Cross-Entropy<br>( <code>from_logits=True</code> )      |
| Multi-class<br>( $c > 2$ ) | $c$ neurons with <code>linear</code><br>activation | Categorical Cross-Entropy<br>( <code>from_logits=True</code> ) |



# Summary: Code examples

## Binary

```
1 model = Sequential([
2     # Skipping the earlier layers
3     Dense(1, activation="sigmoid")
4 ])
5 model.compile(loss="binary_crossentropy")
```

## Binary (logits)

```
1 model = Sequential([
2     # Skipping the earlier layers
3     Dense(1, activation="linear")
4 ])
5 loss = BinaryCrossentropy(from_logits=True)
6 model.compile(loss=loss)
```

## Multi-class

```
1 model = Sequential([
2     # Skipping the earlier layers
3     Dense(n_classes, activation="softmax")
4 ])
5 model.compile(loss="sparse_categorical_crossentropy")
```

## Multi-class (logits)

```
1 model = Sequential([
2     # Skipping the earlier layers
3     Dense(n_classes, activation="linear")
4 ])
5 loss = SparseCategoricalCrossentropy(from_logits=True)
6 model.compile(loss=loss)
```

Both `BinaryCrossentropy` and `SparseCategoricalCrossentropy` live in `keras.losses`.

# Lecture Outline

- Binary Classification
- Multiclass Classification
- **Dense Layers in Matrices**
- Optimisation
- Loss and Derivatives

# Logistic regression

Observations:  $\mathbf{x}_{i,\bullet} \in \mathbb{R}^2$ .

Target:  $y_i \in \{0, 1\}$ .

Predict:  $\hat{y}_i = \mathbb{P}(Y_i = 1)$ .


## The model

For  $\mathbf{x}_{i,\bullet} = (x_{i,1}, x_{i,2})$ :

$$z_i = x_{i,1}w_1 + x_{i,2}w_2 + b$$

$$\hat{y}_i = \sigma(z_i) = \frac{1}{1 + e^{-z_i}}.$$

```
1 x = np.linspace(-10, 10, 100)
2 y = 1/(1 + np.exp(-x))
3 plt.plot(x, y);
```



# Multiple observations

```
1 data = pd.DataFrame({"x_1": [1, 3, 5], "x_2": [2, 4, 6], "y": [0, 1, 1]})
2 data
```

|   | x_1 | x_2 | y |
|---|-----|-----|---|
| 0 | 1   | 2   | 0 |
| 1 | 3   | 4   | 1 |
| 2 | 5   | 6   | 1 |

Let  $w_1 = 1$ ,  $w_2 = 2$  and  $b = -10$ .

```
1 w_1 = 1; w_2 = 2; b = -10
2 data["x_1"] * w_1 + data["x_2"] * w_2 + b
```

```
0    -5
1     1
2     7
dtype: int64
```

# Matrix notation

Have  $\mathbf{X} \in \mathbb{R}^{3 \times 2}$ .

```
1 X_df = data[["x_1", "x_2"]]
2 X = X_df.to_numpy()
3 X
```

```
array([[1, 2],
       [3, 4],
       [5, 6]])
```

Let  $\mathbf{w} = (w_1, w_2)^\top \in \mathbb{R}^{2 \times 1}$ .

```
1 w = np.array([[1], [2]])
2 w
```

```
array([[1],
       [2]])
```

$$\mathbf{z} = \mathbf{X}\mathbf{w} + b, \quad \mathbf{a} = \sigma(\mathbf{z})$$

```
1 z = X.dot(w) + b
2 z
```

```
array([[ -5],
       [  1],
       [  7]])
```

```
1 1 / (1 + np.exp(-z))
```

```
array([[0.01],
       [0.73],
       [1.  ]])
```

# Using a softmax output

Observations:  $\mathbf{x}_{i,\bullet} \in \mathbb{R}^2$ . Predict:  $\hat{y}_{i,j} = \mathbb{P}(Y_i = j)$ .

Target:  $\mathbf{y}_{i,\bullet} \in \{(1, 0), (0, 1)\}$ .

**The model:** For  $\mathbf{x}_{i,\bullet} = (x_{i,1}, x_{i,2})$

$$z_{i,1} = x_{i,1}w_{1,1} + x_{i,2}w_{2,1} + b_1,$$

$$z_{i,2} = x_{i,1}w_{1,2} + x_{i,2}w_{2,2} + b_2.$$

$$\hat{y}_{i,1} = \text{Softmax}_1(\mathbf{z}_i) = \frac{e^{z_{i,1}}}{e^{z_{i,1}} + e^{z_{i,2}}},$$

$$\hat{y}_{i,2} = \text{Softmax}_2(\mathbf{z}_i) = \frac{e^{z_{i,2}}}{e^{z_{i,1}} + e^{z_{i,2}}}.$$

# Multiple observations

1 data

|   | x_1 | x_2 | y_1 | y_2 |
|---|-----|-----|-----|-----|
| 0 | 1   | 2   | 1   | 0   |
| 1 | 3   | 4   | 0   | 1   |
| 2 | 5   | 6   | 0   | 1   |

Choose:

$$w_{1,1} = 1, w_{2,1} = 2,$$

$$w_{1,2} = 3, w_{2,2} = 4, \text{ and}$$

$$b_1 = -10, b_2 = -20.$$

```
1 w_11 = 1; w_21 = 2; b_1 = -10
2 w_12 = 3; w_22 = 4; b_2 = -20
3 data["x_1"] * w_11 + data["x_2"] * w_21 + b_1
```

```
0    -5
1     1
2     7
dtype: int64
```

# Matrix notation

Have  $\mathbf{X} \in \mathbb{R}^{3 \times 2}$ .

```
1 X
```

```
array([[1, 2],
       [3, 4],
       [5, 6]])
```

$\mathbf{W} \in \mathbb{R}^{2 \times 2}$ ,  $\mathbf{b} \in \mathbb{R}^2$

```
1 W = np.array([[1, 3], [2, 4]])
2 b = np.array([-10, -20])
3 display(W); b
```

```
array([[1, 3],
       [2, 4]])
```

```
array([-10, -20])
```

$$\mathbf{Z} = \mathbf{XW} + \mathbf{b}, \quad \mathbf{A} = \text{Softmax}(\mathbf{Z}).$$

```
1 Z = X @ W + b
2 Z
```

```
array([[-5, -9],
       [ 1,  5],
       [ 7, 19]])
```

```
1 np.exp(Z) / np.sum(np.exp(Z),
2   axis=1, keepdims=True)
```

```
array([[9.82e-01, 1.80e-02],
       [1.80e-02, 9.82e-01],
       [6.14e-06, 1.00e+00]])
```




# Lecture Outline

- Binary Classification
- Multiclass Classification
- Dense Layers in Matrices
- **Optimisation**
- Loss and Derivatives

# Gradient-based learning

*In-class demo*

# Gradient descent pitfalls



# Go over all the training data

Called *batch gradient descent*.

```
1 for i in range(num_epochs):  
2     gradient = evaluate_gradient(loss_function, data, weights)  
3     weights = weights - learning_rate * gradient
```

# Pick a random training example

Called *stochastic gradient descent*.

```
1 for i in range(num_epochs):  
2     rnd.shuffle(data)  
3     for example in data:  
4         gradient = evaluate_gradient(loss_function, example, weights)  
5         weights = weights - learning_rate * gradient
```

# Take a group of training examples


Called *mini-batch gradient descent*.

```
1 for i in range(num_epochs):  
2     rnd.shuffle(data)  
3     for b in range(num_batches):  
4         batch = data[b * batch_size : (b + 1) * batch_size]  
5         gradient = evaluate_gradient(loss_function, batch, weights)  
6         weights = weights - learning_rate * gradient
```

# Mini-batch gradient descent


Why?

1. Because we have to (data is too big to shove it all in a single batch)
2. Because it is faster (lots of quick noisy steps *takes longer* than a few slow super accurate steps)
3. The noise helps us jump out of local minima



Noisy gradient means we might jump out of a local minimum.

# Learning rates




Gradient descent with different learning rates

Source: Melissa Renard (2025), adapted from Jeremy Jordan (2018), *Setting the learning rate of your neural network*.




# Learning rates #2



Changing the learning rates for a robot arm.

Source: Matt Henderson (2021), [X post](#)

# Learning rate schedule



Learning curves for various learning rates

In training the learning rate may be tweaked manually.

Source: Melissa Renard (2025), adapted from [Deep Learning for Computer Vision - Stanford Spring 2025](#)

# Lecture Outline

- Binary Classification
- Multiclass Classification
- Dense Layers in Matrices
- Optimisation
- **Loss and Derivatives**

# Example: linear regression

$$\hat{y}(x) = wx + b$$

For some observation  $\{x_i, y_i\}$ , the squared error loss is

$$\text{Loss}_i = (\hat{y}(x_i) - y_i)^2$$

For a batch of the first  $n$  observations the MSE loss is

$$\text{Loss}_{1:n} = \frac{1}{n} \sum_{i=1}^n (\hat{y}(x_i) - y_i)^2$$

# Derivatives

Since  $\hat{y}(x) = wx + b$ ,

$$\frac{\partial \hat{y}(x)}{\partial w} = x \text{ and } \frac{\partial \hat{y}(x)}{\partial b} = 1.$$

As  $\text{Loss}_i = (\hat{y}(x_i) - y_i)^2$ , we know

$$\frac{\partial \text{Loss}_i}{\partial \hat{y}(x_i)} = 2(\hat{y}(x_i) - y_i).$$

# Chain rule

$$\frac{\partial \text{Loss}_i}{\partial \hat{y}(x_i)} = 2(\hat{y}(x_i) - y_i), \quad \frac{\partial \hat{y}(x)}{\partial w} = x, \quad \text{and} \quad \frac{\partial \hat{y}(x)}{\partial b} = 1.$$

Putting this together, we have

$$\frac{\partial \text{Loss}_i}{\partial w} = \frac{\partial \text{Loss}_i}{\partial \hat{y}(x_i)} \times \frac{\partial \hat{y}(x_i)}{\partial w} = 2(\hat{y}(x_i) - y_i) x_i$$

and

$$\frac{\partial \text{Loss}_i}{\partial b} = \frac{\partial \text{Loss}_i}{\partial \hat{y}(x_i)} \times \frac{\partial \hat{y}(x_i)}{\partial b} = 2(\hat{y}(x_i) - y_i).$$

# We need non-zero derivatives

This is why can't use accuracy as the loss function for classification.

Also why we can have the *dead ReLU* problem.

# Stochastic gradient descent (SGD)

Start with  $\theta_0 = (w, b)^\top = (0, 0)^\top$ .

Randomly pick  $i = 5$ , say  $x_i = 5$  and  $y_i = 5$ .

$$\hat{y}(x_i) = 0 \times 5 + 0 = 0 \Rightarrow \text{Loss}_i = (0 - 5)^2 = 25.$$

The partial derivatives are

$$\begin{aligned} \frac{\partial \text{Loss}_i}{\partial w} &= 2(\hat{y}(x_i) - y_i) x_i = 2 \cdot (0 - 5) \cdot 5 = -50, \text{ and} \\ \frac{\partial \text{Loss}_i}{\partial b} &= 2(0 - 5) = -10. \end{aligned}$$

The gradient is  $\nabla \text{Loss}_i = (-50, -10)^\top$ .



# SGD, first iteration

Start with  $\boldsymbol{\theta}_0 = (w, b)^\top = (0, 0)^\top$ .

Randomly pick  $i = 5$ , say  $x_i = 5$  and  $y_i = 5$ .

The gradient is  $\nabla \text{Loss}_i = (-50, -10)^\top$ .

Use learning rate  $\eta = 0.01$  to update

$$\begin{aligned}\boldsymbol{\theta}_1 &= \boldsymbol{\theta}_0 - \eta \nabla \text{Loss}_i \\ &= \begin{pmatrix} 0 \\ 0 \end{pmatrix} - 0.01 \begin{pmatrix} -50 \\ -10 \end{pmatrix} \\ &= \begin{pmatrix} 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 0.5 \\ 0.1 \end{pmatrix} = \begin{pmatrix} 0.5 \\ 0.1 \end{pmatrix}.\end{aligned}$$

# SGD, second iteration

Start with  $\theta_1 = (w, b)^\top = (0.5, 0.1)^\top$ .

Randomly pick  $i = 9$ , say  $x_i = 9$  and  $y_i = 17$ .

The gradient is  $\nabla \text{Loss}_i = (-223.2, -24.8)^\top$ .

Use learning rate  $\eta = 0.01$  to update

$$\begin{aligned}\theta_2 &= \theta_1 - \eta \nabla \text{Loss}_i \\ &= \begin{pmatrix} 0.5 \\ 0.1 \end{pmatrix} - 0.01 \begin{pmatrix} -223.2 \\ -24.8 \end{pmatrix} \\ &= \begin{pmatrix} 0.5 \\ 0.1 \end{pmatrix} + \begin{pmatrix} 2.232 \\ 0.248 \end{pmatrix} = \begin{pmatrix} 2.732 \\ 0.348 \end{pmatrix}.\end{aligned}$$

# Batch gradient descent (BGD)

For the first  $n$  observations  $\text{Loss}_{1:n} = \frac{1}{n} \sum_{i=1}^n \text{Loss}_i$  so

$$\begin{aligned}\frac{\partial \text{Loss}_{1:n}}{\partial w} &= \frac{1}{n} \sum_{i=1}^n \frac{\partial \text{Loss}_i}{\partial w} = \frac{1}{n} \sum_{i=1}^n \frac{\partial \text{Loss}_i}{\hat{y}(x_i)} \frac{\partial \hat{y}(x_i)}{\partial w} \\ &= \frac{1}{n} \sum_{i=1}^n 2(\hat{y}(x_i) - y_i) x_i.\end{aligned}$$

$$\begin{aligned}\frac{\partial \text{Loss}_{1:n}}{\partial b} &= \frac{1}{n} \sum_{i=1}^n \frac{\partial \text{Loss}_i}{\partial b} = \frac{1}{n} \sum_{i=1}^n \frac{\partial \text{Loss}_i}{\hat{y}(x_i)} \frac{\partial \hat{y}(x_i)}{\partial b} \\ &= \frac{1}{n} \sum_{i=1}^n 2(\hat{y}(x_i) - y_i).\end{aligned}$$

# BGD, first iteration ( $\theta_0 = \mathbf{0}$ )

|          | x | y    | y_hat | loss  | dL/dw  | dL/db  |
|----------|---|------|-------|-------|--------|--------|
| <b>0</b> | 1 | 0.99 | 0     | 0.98  | -1.98  | -1.98  |
| <b>1</b> | 2 | 3.00 | 0     | 9.02  | -12.02 | -6.01  |
| <b>2</b> | 3 | 5.01 | 0     | 25.15 | -30.09 | -10.03 |

So  $\nabla \text{Loss}_{1:3}$  is

```
1 nabla = np.array([df["dL/dw"].mean(), df["dL/db"].mean()])
2 nabla
```

```
array([-14.69,  -6.  ])
```

so with  $\eta = 0.1$  then  $\theta_1$  becomes

```
1 theta_1 = theta_0 - 0.1 * nabla
2 theta_1
```

```
array([1.47,  0.6 ])
```

# BGD, second iteration

|          | x | y    | y_hat | loss | dL/dw | dL/db |
|----------|---|------|-------|------|-------|-------|
| <b>0</b> | 1 | 0.99 | 2.07  | 1.17 | 2.16  | 2.16  |
| <b>1</b> | 2 | 3.00 | 3.54  | 0.29 | 2.14  | 1.07  |
| <b>2</b> | 3 | 5.01 | 5.01  | 0.00 | -0.04 | -0.01 |

So  $\nabla \text{Loss}_{1:3}$  is

```
1 nabla = np.array([df["dL/dw"].mean(), df["dL/db"].mean()])
2 nabla
```

array([1.42, 1.07])

so with  $\eta = 0.1$  then  $\theta_2$  becomes

```
1 theta_2 = theta_1 - 0.1 * nabla
2 theta_2
```

array([1.33, 0.49])

# Package Versions

```
1 from watermark import watermark
2 print(watermark(python=True, packages="keras,matplotlib,numpy,pandas,seaborn,scipy,torch")
```

Python implementation: CPython

Python version : 3.14.5

IPython version : 9.13.0

keras : 3.14.1

matplotlib: 3.10.9

numpy : 2.4.4

pandas : 3.0.2

seaborn : 0.13.2

scipy : 1.17.1

torch : 2.11.0

# Recommended viewing

Some very easy-to-follow explanations of these topics, plus catchy tunes:

- Softmax and argmax
- Cross entropy
- Cross entropy derivatives and backpropagation

# Glossary

- accuracy
- classification problem
- confusion matrix
- cross-entropy loss
- metrics
- sigmoid activation function
- softmax activation
- batch gradient descent
- batches, batch size
- global minimum, local minimum
- gradient-based learning, hill-climbing
- learning rate, learning rate schedule
- plateau
- stochastic gradient descent
- mini-batch gradient descent