

Entity Embedding

ACTL3143 & ACTL5111 Deep Learning for Actuaries
Patrick Laub

Lecture Outline

- **Entity Embedding**
- The French Motor Dataset
- From One-Hot Encoding to Embeddings
- Keras' Functional API
- French Motor Dataset with Embeddings
- Scale by Exposure

Each category used to be a column

When we preprocessed categorical variables, *one-hot encoding* gave a *column for every category*, so you could always read off what each number meant:

Category	One-hot vector
R11	$[1, 0, 0, \dots, 0]$
R24	$[0, 1, 0, \dots, 0]$
R93	$[0, 0, \dots, 1, 0]$

Each column *means one specific category*, but the vector is *sparse* (mostly zeros) and *long* (one entry per category), and every category sits the *same distance* from every other.

Entity embeddings: short, dense & learned

An *entity embedding* (Guo & Berkhahn, 2016) replaces that sparse column-per-category block with a *short, dense vector* that the network *learns during training*:

$$R11 \longrightarrow [0.07, -0.42], \quad R24 \longrightarrow [0.11, -0.39]$$

- the vector is *short* — e.g. 2 numbers instead of 22 columns;
- the dimensions are *learned*, not categories — there is no “R11 column”;
- categories with a *similar effect on the target* tend to cluster together.

This is particularly useful when the categorical variable can take on a large number of different levels (it has high cardinality). Though, in that situation, *rare* levels are fit from very little data, so they may not be placed optimally.

It's the same idea as word embeddings

- A *word embedding* is just an entity embedding where the “entity” is a *word*.
- *Entity embedding* applies the same trick to *any* categorical variable: region, vehicle brand, occupation, postcode, ...
- Typically word embeddings are *pre-trained* on huge text corpora — i.e. it is a *transfer learning* technique — though here we have labels so learn the embeddings as part of the supervised task.

They are both, in essence, a *lookup table of learned vectors*.

Lecture Outline

- Entity Embedding
- **The French Motor Dataset**
- From One-Hot Encoding to Embeddings
- Keras' Functional API
- French Motor Dataset with Embeddings
- Scale by Exposure

Imports needed for these demos

```
1 import random
2 from pathlib import Path
3
4 import matplotlib.pyplot as plt
5 import numpy as np
6 import pandas as pd
7
8 from keras.models import Sequential, Model
9 from keras.layers import Dense, Input, Embedding, Reshape, Concatenate
10 from keras.callbacks import EarlyStopping
11 from keras.utils import plot_model
12
13 from sklearn import set_config
14 from sklearn.datasets import fetch_openml
15 from sklearn.model_selection import train_test_split
16 from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, StandardScaler
17 from sklearn.compose import make_column_transformer
18
19 set_config(transform_output="pandas")
```

Revisit the French motor dataset

	IDpol	ClaimNb	Exposure	Area	VehPower	VehAge	DrivAge	BonI
0	1.0	1.0	0.10000	D	5.0	0.0	55.0	50.0
1	3.0	1.0	0.77000	D	5.0	0.0	55.0	50.0
...	...	...	...	...	...	...	...	...
678011	6114329.0	0.0	0.00274	B	4.0	0.0	60.0	50.0
678012	6114330.0	0.0	0.00274	B	7.0	6.0	29.0	54.0

678013 rows × 12 columns

Source: Noll et al. (2020).

Data dictionary

Variable	Description	Preprocessing
IDpol	Policy number (unique identifier)	Dropped
ClaimNb	Number of claims on the given policy	Target
Exposure*	Total exposure in yearly units	Normalised
Area	Area code (ordinal)	Ordinal Encode
VehPower	Power of the car (ordinal encoded)	Normalised
VehAge	Age of the car in years	Normalised
DrivAge	Age of the (most common) driver in years	Normalised
BonusMalus	Bonus–malus level between 50 and 230 (with reference level 100)	Normalised
VehBrand*	Car brand (nominal)	One-hot
VehGas	Diesel or regular fuel car (binary)	One-hot
Density	Density of inhabitants per km ² in the city of the living place of the driver	Normalised
Region*	Regions in France (prior to 2016)	One-hot

* The three variables we single out later: *Region* & *VehBrand* get embeddings, and *Exposure* becomes an offset.

The model

Have $\{(\mathbf{x}_i, y_i)\}_{i=1, \dots, n}$ for $\mathbf{x}_i \in \mathbb{R}^{47}$ and $y_i \in \mathbb{N}_0$.

Assume the distribution

$$Y_i \sim \text{Poisson}(\lambda(\mathbf{x}_i))$$

We have $\mathbb{E}Y_i = \lambda(\mathbf{x}_i)$. The NN takes $\mathbf{x}_i$ & predicts $\mathbb{E}Y_i$.

i Note

For insurance, *this is a bit weird*. The exposures are different for each policy.

$\lambda(\mathbf{x}_i)$ is the expected number of claims for the duration of policy i 's contract.

Normally, $\text{Exposure}_i \notin \mathbf{x}_i$, and $\lambda(\mathbf{x}_i)$ is the expected rate *per year*, then

$$Y_i \sim \text{Poisson}(\text{Exposure}_i \times \lambda(\mathbf{x}_i)).$$

What values do we see in the data?

```
1 X_train_raw["Area"].value_counts()
2 X_train_raw["VehBrand"].value_counts()
3 X_train_raw["VehGas"].value_counts()
4 X_train_raw["Region"].value_counts()
```

Area

C 5514
D 4116

...

B 2387
F 444

Name: count, Length: 6, dtype: int64

VehGas

Regular 10658
Diesel 8092

Name: count, dtype: int64

VehBrand

B1 4998
B2 4906

...

B11 283
B14 140

Name: count, Length: 11, dtype: int64

Region

R24 6493
R82 2112

...

R42 48
R43 26

Name: count, Length: 22, dtype: int64

How we preprocessed last time

```

1 ct = make_column_transformer(
2     (OneHotEncoder(sparse_output=False, drop="first"), ["VehGas", "VehBrand", "Region"]),
3     (OrdinalEncoder(), ["Area"]),
4     remainder=StandardScaler(),
5     verbose_feature_names_out=False
6 )
7 X_train = ct.fit_transform(X_train_raw)

```

```
1 X_train_raw.head(3)
```

	Exposure	Area	VehPower	VehAge
0	1.00	A	7.0	8.0
1	0.79	B	7.0	7.0
2	1.00	C	6.0	13.0

```
1 X_train.head(3)
```


	VehGas_Regular	VehBrand_B10	Ve
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	1.0	0.0	0.0

3 rows × 39 columns

Lecture Outline

- Entity Embedding
- The French Motor Dataset
- **From One-Hot Encoding to Embeddings**
- Keras' Functional API
- French Motor Dataset with Embeddings
- Scale by Exposure

Region column



French Administrative Regions

Source: [Wikimedia](#)

One-hot encoding

```

1 oh = OneHotEncoder(sparse_output=False)
2 X_train_oh = oh.fit_transform(X_train_raw[["Region"]])
3 X_test_oh = oh.transform(X_test_raw[["Region"]])
4 print(list(X_train_raw["Region"][:5]))
5 X_train_oh.head()

```

['R24', 'R21', 'R53', 'R24', 'R82']

	Region_R11	Region_R21	Region_R22	Region_R23	Region_R24	Region_R25
0	0.0	0.0	0.0	0.0	1.0	0.0
1	0.0	1.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...
3	0.0	0.0	0.0	0.0	1.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 22 columns

Train on one-hot inputs

```
1 num_regions = len(oh.categories_[0])
2
3 random.seed(12)
4 model = Sequential([
5     Input(shape=(num_regions,)),
6     Dense(2),
7     Dense(1, activation="exponential")
8 ])
9
10 model.compile(optimizer="adam", loss="poisson")
11
12 es = EarlyStopping(verbose=True)
13 hist = model.fit(X_train_oh, y_train, epochs=100, verbose=0, validation_split=0.2, callba
14
15 hist.history["val_loss"][-1]
```

Epoch 6: early stopping

0.7679625153541565

Make a fake batch of data

```
1 X = np.eye(num_regions)
2 pd.DataFrame(X, columns=oh.categories_[0])
```

	R11	R21	R22	R23	R24	R25	R26
0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	1.0	0.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...
20	0.0	0.0	0.0	0.0	0.0	0.0	0.0
21	0.0	0.0	0.0	0.0	0.0	0.0	0.0

22 rows × 22 columns

```
1 model.layers[0](X)
```

```
tensor([[ -0.0646, -0.2346],
        [ 0.1463, -0.0323],
        [-0.3050, -0.2913],
        [ 0.0257, -0.3189],
        [ 0.3848, -0.2778],
        [ 0.0983, -0.4137],
        [ 0.2235, -0.2502],
        [ 0.1192,  0.3202],
        [ 0.6295, -0.4385],
        [ 0.2008, -0.0306],
        [-0.2030, -0.3797],
        [ 0.7612, -0.2256],
        [ 0.5068, -0.0436],
        [ 0.7867, -0.4326],
        [-0.2026, -0.1788],
        [-0.2430, -0.0941],
```

The first layer

```
1 layer = model.layers[0]
2 W, b = layer.get_weights()
3 X.shape, W.shape, b.shape
```

```
((22, 22), (22, 2), (2,))
```

```
1 X @ W + b
```

```
array([[ -0.06, -0.23],
       [ 0.15, -0.03],
       [-0.3 , -0.29],
       [ 0.03, -0.32],
       [ 0.38, -0.28],
       [ 0.1 , -0.41],
       [ 0.22, -0.25],
       [ 0.12,  0.32],
       [ 0.63, -0.44],
       [ 0.2 , -0.03],
       [-0.2 , -0.38],
       [ 0.76, -0.23],
       [ 0.51, -0.04],
       [ 0.79, -0.43],
       [-0.2 , -0.18],
       [-0.24, -0.09],
```

```
1 W + b
```

```
array([[ -0.06, -0.23],
       [ 0.15, -0.03],
       [-0.3 , -0.29],
       [ 0.03, -0.32],
       [ 0.38, -0.28],
       [ 0.1 , -0.41],
       [ 0.22, -0.25],
       [ 0.12,  0.32],
       [ 0.63, -0.44],
       [ 0.2 , -0.03],
       [-0.2 , -0.38],
       [ 0.76, -0.23],
       [ 0.51, -0.04],
       [ 0.79, -0.43],
       [-0.2 , -0.18],
       [-0.24, -0.09],
```

Just a look-up operation

```
1 display(list(oh.categories_[0]))
```

```
['R11',
 'R21',
 'R22',
 'R23',
 'R24',
 'R25',
 'R26',
 'R31',
 'R41',
 'R42',
 'R43',
 'R52',
 'R53',
 'R54',
 'R72',
 'R73',
```

```
1 W + b
```

```
array([[ -0.06, -0.23],
       [ 0.15, -0.03],
       [-0.3 , -0.29],
       [ 0.03, -0.32],
       [ 0.38, -0.28],
       [ 0.1 , -0.41],
       [ 0.22, -0.25],
       [ 0.12,  0.32],
       [ 0.63, -0.44],
       [ 0.2 , -0.03],
       [-0.2 , -0.38],
       [ 0.76, -0.23],
       [ 0.51, -0.04],
       [ 0.79, -0.43],
       [-0.2 , -0.18],
       [-0.24, -0.09],
```

Each category maps to one row of $W + b$ — that row *is* its embedding. No matrix multiply needed.

Turn the region into an index

```
1 oe = OrdinalEncoder()
2 X_train_reg = oe.fit_transform(X_train_raw[["Region"]])
3 X_test_reg = oe.transform(X_test_raw[["Region"]])
4
5 for i, reg in enumerate(oe.categories_[0][:3]):
6     print(f"The Region value {reg} gets turned into {i}.")
```

The Region value R11 gets turned into 0.

The Region value R21 gets turned into 1.

The Region value R22 gets turned into 2.

Use an Embedding layer

```

1 num_regions = X_train_raw["Region"].unique()
2
3 random.seed(12)
4 model = Sequential([
5     Input(shape=(1,)),
6     Embedding(input_dim=num_regions, output_dim=2),
7     Dense(1, activation="exponential")
8 ])
9
10 model.compile(optimizer="adam", loss="poisson")

```

```

1 es = EarlyStopping(verbose=True)
2 hist = model.fit(X_train_reg, y_train, epochs=100, verbose=0,
3     validation_split=0.2, callbacks=[es])
4 hist.history["val_loss"][-1]

```

Epoch 4: early stopping

0.7677991390228271

```
1 model.layers
```

[<Embedding name=embedding, built=True>, <Dense name=dense_2, built=True>]

Aside: the output here is shaped `(None, 1, 1)`, not `(None, 1)` — a harmless extra axis from the `Embedding` that we will tidy up later

Keras' Embedding Layer

```
1 model.layers[0].get_weights()[0]
```

```
array([[ 0.05, -0.07],
       [-0.02,  0.03],
       [-0.04, -0.02],
       [ 0.09, -0.12],
       [ 0.24, -0.22],
       [ 0.16, -0.19],
       [ 0.17, -0.17],
       [-0.05,  0.09],
       [ 0.36, -0.33],
       [ 0.03, -0.01],
       [ 0.02, -0.07],
       [ 0.36, -0.3 ],
       [ 0.21, -0.16],
       [ 0.42, -0.37],
       [-0.02, -0.02],
       [-0.06,  0.02],
```

```
1 X_train_raw["Region"].head(4)
```

```
0    R24
1    R21
2    R53
3    R24
Name: Region, dtype: str
```

```
1 X_sample = X_train_reg[:4].to_numpy()
2 X_sample
```

```
array([[ 4.],
       [ 1.],
       [12.],
       [ 4.]])
```

```
1 enc_tensor = model.layers[0](X_sample)
2 keras.ops.convert_to_numpy(enc_tensor).sc
```

```
array([[ 0.24, -0.22],
       [-0.02,  0.03],
       [ 0.21, -0.16],
       [ 0.24, -0.22]], dtype=float32)
```

The embedding dimension

- Entity embedding is especially useful when there is a very large number of categories (such as 1,000 or 10,000) and you want to reduce the number of columns.
- The embedding dimension is really a *hyperparameter to tune* (e.g. by validation loss); there is no settled formula for it.
- One heuristic is to use $n^{1/4}$ for a variable with n categories (The TensorFlow Team, 2017), but there's no theoretical justification, and others disagree.
- In this case, a nice side benefit is that 2-dimensional embeddings can be plotted directly.

The learned embeddings

```

1 points = model.layers[0].get_weights()[0]
2 plt.figure(figsize=(3, 3))
3 plt.scatter(points[:,0], points[:,1])
4 for i in range(num_regions):
5     plt.text(points[i,0]+0.01, points[i,1]

```






FIG. 3. The learned German state embedding is mapped to a 2D space with t-SNE. The relative positions of German states here resemble that on the real German map surprisingly well.

Reproduced from the Guo & Berkhahn (2016, p. 7).

Embeddings improve during training



Embeddings move during training to wherever they need to be to minimise training loss.

Source: Marcus Lautier (2022).

Embeddings & other inputs




Illustration of a neural network with both continuous and categorical inputs.

We can't do this with Sequential models...

Source: LotusLabs Blog, [Accurate insurance claims prediction with Deep Learning](#).

Lecture Outline

- Entity Embedding
- The French Motor Dataset
- From One-Hot Encoding to Embeddings
- **Keras' Functional API**
- French Motor Dataset with Embeddings
- Scale by Exposure

Converting Sequential models

```

1 random.seed(12)
2
3 model = Sequential([
4     Input(shape=(X_train_oh.shape[1],)),
5     Dense(30, "leaky_relu"),
6     Dense(1, "exponential")
7 ])
8
9 model.compile(
10     optimizer="adam",
11     loss="poisson")
12
13 hist = model.fit(
14     X_train_oh, y_train,
15     epochs=1, verbose=0,
16     validation_split=0.2)
17 hist.history["val_loss"][-1]

```

0.7695862650871277


```

1 random.seed(12)
2
3 inputs = Input(shape=(X_train_oh.shape[1]
4 x = Dense(30, "leaky_relu")(inputs)
5 out = Dense(1, "exponential")(x)
6 model = Model(inputs, out)
7
8 model.compile(
9     optimizer="adam",
10     loss="poisson")
11
12 hist = model.fit(
13     X_train_oh, y_train,
14     epochs=1, verbose=0,
15     validation_split=0.2)
16 hist.history["val_loss"][-1]

```

0.7695212364196777

Wide & Deep network



An illustration of the wide & deep network architecture.

Add a *skip connection* from input to output layers.

```

1 inp = Input(shape=X_train.shape[1:])
2 hidden1 = Dense(30, "leaky_relu")(inp)
3 hidden2 = Dense(30, "leaky_relu")(hidden1)
4 concat = Concatenate()(
5     [inp, hidden2])
6 output = Dense(1)(concat)
7 model = Model(inputs=[inp], outputs=[output])
  
```


Naming the layers

For complex networks, it is often useful to give meaningful names to the layers.

```
1 input_ = Input(shape=X_train.shape[1:], name="input")
2 hidden1 = Dense(30, activation="leaky_relu", name="hidden1")(input_)
3 hidden2 = Dense(30, activation="leaky_relu", name="hidden2")(hidden1)
4 concat = Concatenate(name="combined")([input_, hidden2])
5 output = Dense(1, name="output")(concat)
6 model = Model(inputs=[input_], outputs=[output])
```

Inspecting a complex model

```
1 plot_model(model, show
```



```
1 model.summary(line_length=60)
```

Model: "functional_5"

Layer (type)	Output Shape	Param #	Connected to
input (InputLayer)	(None, 39)	0	–
hidden1 (Dense)	(None, 30)	1,200	input[0][0]
hidden2 (Dense)	(None, 30)	930	hidden1[0][0]
combined (Concatenate)	(None, 69)	0	input[0][0], hidden2[0][0]
output (Dense)	(None, 1)	70	combined[0][...]

Total params: 2,200 (8.59 KB)
 Trainable params: 2,200 (8.59 KB)
 Non-trainable params: 0 (0.00 B)

Lecture Outline

- Entity Embedding
- The French Motor Dataset
- From One-Hot Encoding to Embeddings
- Keras' Functional API
- **French Motor Dataset with Embeddings**
- Scale by Exposure

The desired architecture




Illustration of a neural network with both continuous and categorical inputs.

Source: LotusLabs Blog, [Accurate insurance claims prediction with Deep Learning](#).

Preprocess all French motor inputs

Transform the categorical variables to integers:

```
1 num_brands, num_regions = X_train_raw[["VehBrand", "Region"]].nunique()
2
3 ct = make_column_transformer(
4     (OrdinalEncoder(), ["VehBrand", "Region", "Area", "VehGas"]),
5     remainder=StandardScaler(),
6     verbose_feature_names_out=False
7 )
8 X_train = ct.fit_transform(X_train_raw)
9 X_test = ct.transform(X_test_raw)
```

Split the brand and region data apart from the rest:

```
1 X_train_brand = X_train["VehBrand"]
2 X_train_region = X_train["Region"]
3 X_train_rest = X_train.drop(["VehBrand", "Region"], axis=1)
4
5 X_test_brand = X_test["VehBrand"]
6 X_test_region = X_test["Region"]
7 X_test_rest = X_test.drop(["VehBrand", "Region"], axis=1)
```

Organise the inputs

Make a Keras `Input` for: vehicle brand, region, & others.

```
1 veh_brand = Input(shape=(1,), name="veh_brand")
2 region = Input(shape=(1,), name="region")
3 other_inputs = Input(shape=X_train_rest.shape[1:], name="other_inputs")
```

Create embeddings and join them with the other inputs.

```
1 random.seed(1337)
2 veh_brand_ee = Embedding(input_dim=num_brands, output_dim=2,
3     name="veh_brand_ee")(veh_brand)
4 veh_brand_ee = Reshape(target_shape=(2,))(veh_brand_ee)
5
6 region_ee = Embedding(input_dim=num_regions, output_dim=2, name="region_ee")(region)
7 region_ee = Reshape(target_shape=(2,))(region_ee)
8
9 x = Concatenate(name="combined")([veh_brand_ee, region_ee, other_inputs])
```

Complete the model and fit it


Feed the combined embeddings & continuous inputs to some normal dense layers.

```
1 x = Dense(30, "relu", name="hidden")(x)
2 out = Dense(1, "exponential", name="out")(x)
3
4 model = Model([veh_brand, region, other_inputs], out)
5 model.compile(optimizer="adam", loss="poisson")
6
7 hist = model.fit((X_train_brand, X_train_region, X_train_rest),
8                 y_train, epochs=100, verbose=0,
9                 callbacks=[EarlyStopping(patience=5)], validation_split=0.2)
10 np.min(hist.history["val_loss"])
```

```
np.float64(0.6855398416519165)
```


Plotting this model

```
1 plot_model(model, show_layer_names=True)
```



Why we need to reshape

```
1 plot_model(model, show_layer_names=True, show_shapes=True)
```



Embedding turns (None, 1) into (None, 1, 2); Reshape drops the extra axis to (None, 2) so we can Concatenate.

Lecture Outline

- Entity Embedding
- The French Motor Dataset
- From One-Hot Encoding to Embeddings
- Keras' Functional API
- French Motor Dataset with Embeddings
- **Scale by Exposure**

Two different models

Have $\{(\mathbf{x}_i, y_i)\}_{i=1, \dots, n}$ for $\mathbf{x}_i \in \mathbb{R}^{47}$ and $y_i \in \mathbb{N}_0$.

Model 1: Say $Y_i \sim \text{Poisson}(\lambda(\mathbf{x}_i))$.

But, the exposures are different for each policy. $\lambda(\mathbf{x}_i)$ is the expected number of claims for the duration of policy i 's contract.

Model 2: Say $Y_i \sim \text{Poisson}(\text{Exposure}_i \times \lambda(\mathbf{x}_i))$.

Now, $\text{Exposure}_i \notin \mathbf{x}_i$, and $\lambda(\mathbf{x}_i)$ is the rate *per year*.

Just take continuous variables

```
1 ct = make_column_transformer(  
2     ("passthrough", ["Exposure"]),  
3     ("drop", ["VehBrand", "Region", "Area", "VehGas"]),  
4     remainder=StandardScaler(),  
5     verbose_feature_names_out=False  
6 )  
7 X_train = ct.fit_transform(X_train_raw)  
8 X_test = ct.transform(X_test_raw)
```

Split exposure apart from the rest:

```
1 X_train_exp = X_train["Exposure"]  
2 X_test_exp = X_test["Exposure"]  
3 X_train_rest = X_train.drop("Exposure", axis=1)  
4 X_test_rest = X_test.drop("Exposure", axis=1)
```

Organise the inputs:

```
1 exposure = Input(shape=(1,), name="exposure")  
2 other_inputs = Input(shape=X_train_rest.shape[1:], name="other_inputs")
```

Make & fit the model

Feed the continuous inputs to some normal dense layers.

```

1 random.seed(1337)
2 x = Dense(30, "relu", name="hidden1")(other_inputs)
3 x = Dense(30, "relu", name="hidden2")(x)
4 lambda_ = Dense(1, "exponential", name="lambda")(x)

1 out = lambda_ * exposure
2 model = Model([exposure, other_inputs], out)
3 model.compile(optimizer="adam", loss="poisson")
4
5 es = EarlyStopping(patience=10, restore_best_weights=True, verbose=1)
6 hist = model.fit((X_train_exp, X_train_rest), y_train, epochs=100, verbose=0,
7                 callbacks=[es], validation_split=0.2)
8 np.min(hist.history["val_loss"])

```


Epoch 29: early stopping

Restoring model weights from the end of the best epoch: 19.

np.float64(0.9100556373596191)

Plot the model

```
1 plot_model(model, show_layer_names=True)
```



Further reading

Avanzi et al. (2024) focus on how to handle the case of high-cardinality categorical features, and propose a new architecture specifically tailored for actuarial purposes.

Package Versions

```
1 from watermark import watermark
2 print(watermark(python=True, packages="keras,matplotlib,numpy,pandas,seaborn,scipy,torch")
```

Python implementation: CPython

Python version : 3.14.5

IPython version : 9.15.0

keras : 3.15.0

matplotlib: 3.11.0

numpy : 2.5.0

pandas : 3.0.3

seaborn : 0.13.2

scipy : 1.18.0

torch : 2.12.1

Glossary

- entity embeddings
- Input layer
- Keras functional API
- Reshape layer
- skip connection
- wide & deep network

References

- Avanzi, B., Taylor, G., Wang, M., & Wong, B. (2024). Machine learning with high-cardinality categorical features in actuarial applications. *ASTIN Bulletin: The Journal of the IAA*, 54(2), 213–238. <https://doi.org/10.1017/asb.2024.7>
- Guo, C., & Berkahn, F. (2016). Entity embeddings of categorical variables. *arXiv Preprint arXiv:1604.06737*.
- Noll, A., Salzmänn, R., & Wüthrich, M. V. (2020). Case study: French motor third-party liability claims. *SSRN Manuscript* 3164764.
- The TensorFlow Team. (2017). *Introducing TensorFlow feature columns*. Google Developers Blog. <https://developers.googleblog.com/introducing-tensorflow-feature-columns/>