

Time Series & Recurrent Neural Networks

ACTL3143 & ACTL5111 Deep Learning for Actuaries
Patrick Laub

Lecture Outline

- **Time Series**
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series

Tabular data vs time series data

Tabular data

We have a dataset $\{\mathbf{x}_i, y_i\}_{i=1}^n$ which we assume are i.i.d. observations.

Brand	Mileage	# Claims
BMW	101 km	1
Audi	432 km	0
Volvo	3 km	5
$\vdots$	$\vdots$	$\vdots$

The goal is to *predict* the y for some covariates $\mathbf{x}$.

Time series data

Have a sequence $\{\mathbf{x}_t, y_t\}_{t=1}^T$ of observations taken at regular time intervals.

Date	Humidity	Temp.
Jan 1	60%	20 °C
Jan 2	65%	22 °C
Jan 3	70%	21 °C
$\vdots$	$\vdots$	$\vdots$

The task is to *forecast* future values based on the past.

Attributes of time series data

- **Temporal ordering:** The order of the observations matters.
- **Trend:** The general direction of the data.
- **Noise:** Random fluctuations in the data.
- **Seasonality:** Patterns that repeat at regular intervals.

Question

What will be the temperature in Berlin tomorrow? What information would you use to make a prediction?

Australian financial stocks

```

1  # First, install yfinance if you haven't already:
2  # pip install yfinance
3  import yfinance as yf
4  import pandas as pd
5  from datetime import date
6
7  # 1. Define the ASX tickers (Yahoo uses ".AX" for Australian stocks
8  #    and "^AXJO" for the S&P/ASX 200 index)
9  tickers = {
10     "ANZ":    "ANZ.AX",
11     "BOQ":    "BOQ.AX",
12     "CBA":    "CBA.AX",
13     "NAB":    "NAB.AX",
14     "QBE":    "QBE.AX",
15     "SUN":    "SUN.AX",
16     "WBC":    "WBC.AX",
17     "ASX200": "^AXJO"
18 }
19
20 # 2. Choose your date range
21 start = "1999-01-01"
22 end   = date.today().isoformat() # e.g. "2025-06-29"
23
24 # 3. Download the data
25 #    This returns a Panel-like DataFrame: columns are tickers, rows are trading dates.
26 raw_data = yf.download(
27     tickers=list(tickers.values()),
28     start=start,
29     end=end,
30     progress=False
31 )
32
33 raw_data.to_csv("asx_raw_data.csv") # Optional: Save raw data for inspection
34

```

Australian financial stocks


```
1 stocks = pd.read_csv("data/aus_fin_stocks.csv")
2 stocks
```

	Date	ANZ	ASX200	BOQ	CBA	NAB	QBE	SUN	WBC
0	1999-01-01	8.413491	NaN	NaN	14.560169	NaN	NaN	7.965791	NaN
1	1999-01-04	8.476515	2732.199951	NaN	14.402927	12.995510	2.648447	7.965791	6.370629
2	1999-01-05	8.452881	2716.600098	NaN	14.409224	13.169948	2.564247	7.965791	6.485921
...	...	...	...	...	...	...	...	...	...
6742	2025-06-25	29.100000	8559.200195	7.86	191.399994	40.049999	23.480000	21.750000	34.540001
6743	2025-06-26	29.740000	8550.799805	7.86	190.710007	39.889999	23.330000	21.459999	34.570000
6744	2025-06-27	29.200001	8514.200195	7.77	185.360001	39.259998	23.219999	21.330000	33.900002

6745 rows × 9 columns

Plot

```
1 stocks.plot()
```



Data types and NA values

```
1 stocks.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 6745 entries, 0 to 6744
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        6745 non-null   str
1   ANZ         6745 non-null   float64
2   ASX200      6691 non-null   float64
3   BOQ         6599 non-null   float64
4   CBA         6745 non-null   float64
5   NAB         6744 non-null   float64
6   QBE         6744 non-null   float64
7   SUN         6745 non-null   float64
8   WBC         6744 non-null   float64
dtypes: float64(8), str(1)
memory usage: 540.3 KB
```

```
1 asx200 = stocks.pop("ASX200")
```

```
1 for col in stocks.columns:
2     print(f"{col}: {stocks[col].isna().sum()}")
```

```
Date: 0
ANZ: 0
ASX200: 54
BOQ: 146
CBA: 0
NAB: 1
QBE: 1
SUN: 0
WBC: 1
```

Set the index to the date

```
1 stocks["Date"] = pd.to_datetime(stocks["Date"])
2 stocks = stocks.set_index("Date") # or `stocks.set_index("Date", inplace=True)`
3 stocks
```

	ANZ	BOQ	CBA	NAB	QBE	SUN	WBC
Date							
1999-01-01	8.413491	NaN	14.560169	NaN	NaN	7.965791	NaN
1999-01-04	8.476515	NaN	14.402927	12.995510	2.648447	7.965791	6.370629
1999-01-05	8.452881	NaN	14.409224	13.169948	2.564247	7.965791	6.485921
...	...	...	...	...	...	...	...
2025-06-25	29.100000	7.86	191.399994	40.049999	23.480000	21.750000	34.540001
2025-06-26	29.740000	7.86	190.710007	39.889999	23.330000	21.459999	34.570000
2025-06-27	29.200001	7.77	185.360001	39.259998	23.219999	21.330000	33.900002


6745 rows × 7 columns

Plot II

```

1 stocks.plot()
2 plt.ylabel("Stock Price ($)")
3 plt.legend(loc="upper center", bbox_to_anchor=(0.5, -0.5), ncol=4);

```



Can index using dates I

```
1 stocks.loc["2010-1-4":"2010-01-8"]
```

	ANZ	BOQ	CBA	NAB	QBE	SUN	WBC
Date							
2010-01-04	18.874071	8.412694	34.510429	14.440043	13.831507	10.478957	14.958415
2010-01-05	18.964771	8.520640	35.032455	14.624498	13.902833	10.636262	15.082577
2010-01-06	18.684425	8.477461	35.208561	14.339908	13.688861	10.406354	15.011630
2010-01-07	18.239164	8.218388	34.868923	14.223969	13.441965	10.551559	14.810603
2010-01-08	18.346357	8.419890	35.321777	14.176538	13.590102	10.612061	14.869730

Note, these ranges are *inclusive*, not like Python's normal slicing.

Can index using dates II

So to get 2019's December and all of 2020 for CBA:


```
1 stocks.loc["2019-12":"2020", ["CBA"]]
```

CBA	
Date	
2019-12-02	66.193298
2019-12-03	64.494362
2019-12-04	63.250660
...	...
2020-12-29	70.822792
2020-12-30	70.468712
2020-12-31	69.221031

275 rows × 1 columns


Can look at the first differences

```
1 stocks.diff().plot()
2 plt.ylabel("Daily Price Changes ($)")
3 plt.legend(loc="upper center", bbox_to_anchor=(0.5, -0.5), ncol=4);
```



Can look at the percentage changes

```
1 stocks.pct_change().plot()  
2 plt.ylabel("Daily Returns (%)")  
3 plt.legend(loc="upper center", bbox_to_anchor=(0.5, -0.5), ncol=4);
```




Focus on one stock

```
1 stock = stocks[["CBA"]].copy()
2 stock
```

CBA	
Date	
1999-01-01	14.560169
1999-01-04	14.402927
1999-01-05	14.409224
...	...
2025-06-25	191.399994
2025-06-26	190.710007
2025-06-27	185.360001

6745 rows × 1 columns

```
1 stock.plot()
2 plt.ylabel("Stock Price ($)");
```



```
1 stock.isna().sum()
```



```
CBA      0
dtype: int64
```

Convert to log returns

Instead of working with raw prices, we'll work with daily log returns:

```
1 # Calculate log returns
2 stock_log = np.log(stock / stock.shift(1))
3 stock_log.head()
```

	CBA
Date	
1999-01-04	-0.010858
1999-01-05	0.000437
1999-01-06	0.008042
1999-01-07	0.025859
1999-01-08	-0.021323



Fill in the missing values

```
1 asx200 = pd.DataFrame(asx200).set_index(stocks.index)
2 missing_day = asx200.index[asx200["ASX200"].isna()][1]
3 prev_day = missing_day - pd.Timedelta(days=1)
4 after = missing_day + pd.Timedelta(days=3)
```

```
1 asx200.loc[prev_day:after]
```

```
1 asx200 = asx200.ffill()
2 asx200.loc[prev_day:after]
```

ASX200	
Date	
1999-01-25	2713.100098
1999-01-26	NaN
1999-01-27	2738.800049
1999-01-28	2765.100098
1999-01-29	2781.699951

ASX200	
Date	
1999-01-25	2713.100098
1999-01-26	2713.100098
1999-01-27	2738.800049
1999-01-28	2765.100098
1999-01-29	2781.699951

Lecture Outline

- Time Series
- **Baseline forecasts**
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series


Helper functions

```
1 def log_to_price(log_returns, initial_price):
2     """Convert log returns to raw prices given an initial price."""
3     # Use cumulative sum of log returns for numerical stability
4     #  $P_t = P_0 * \exp(\text{sum of log returns from 1 to } t)$ 
5     cumulative_log_returns = log_returns.cumsum()
6     prices = initial_price * np.exp(cumulative_log_returns)
7
8     return prices
9
10 def get_last_price(stock_df, cutoff_date):
11     """Get the last known price before the forecast period starts."""
12     last_known_date = stock_df.loc[:cutoff_date].index[-1]
13     return stock_df.loc[last_known_date, "CBA"]
```

Persistence forecast

Predict the next value to be the same as the current value.

```
1 last_price = get_last_price(stock, cutoff_date="2024-12")
2 log_persistence = pd.Series(0.0, index=stock_log.loc["2025":].index)
3 persistence_prices = log_to_price(log_persistence, last_price)
```




Extrapolate the trend

```
1 recent_log_returns = stock_log.loc["2022-01":"2024-12"]
2 trend_log = recent_log_returns.mean().values[0]
3 print(f"Average daily log return: {trend_log:.6f}")
```


Average daily log return: 0.000709

```
1 log_trend = pd.Series(trend_log, index=stock_log.loc["2025:"].index)
2 trend_prices = log_to_price(log_trend, last_price)
```




Trend fitted

```
1 # Create trend forecast for the recent period to show fitted trend
2 recent_trend_log = pd.Series(trend_log, index=recent_log_returns.index)
3 trend_start_price = get_last_price(stock, cutoff_date=recent_log_returns.index[0].strftime('%Y-%m-%d'))
4 recent_trend_prices = log_to_price(recent_trend_log, trend_start_price)
```



Trend forecasts



Which is better?

If we look at the mean squared error (MSE) of the two models:

```
1 # Calculate MSE using the actual forecasts we computed
2 actual_prices = stock.loc["2025":, "CBA"]
3 persistence_mse = mean_squared_error(actual_prices, persistence_prices)
4 trend_mse = mean_squared_error(actual_prices, trend_prices)
5 persistence_mse, trend_mse
```

```
(254.04075100411256, 99.28717915684173)
```

Use the history

```
1 cba_log_shifted = stock_log["CBA"].head().shift(1)
2 both = pd.concat([stock_log["CBA"].head(), cba_log_shifted], axis=1, keys=["Today", "Yesterday"])
3 both
```

	Today	Yesterday
Date		
1999-01-04	-0.010858	NaN
1999-01-05	0.000437	-0.010858
1999-01-06	0.008042	0.000437
1999-01-07	0.025859	0.008042
1999-01-08	-0.021323	0.025859

```
1 def lagged_timeseries(df, target, window=40):
2     lagged = pd.DataFrame()
3     for i in range(window, 0, -1):
4         lagged[f"T-{i}"] = df[target].shift(i)
5     lagged["T"] = df[target].values
6     return lagged
```

Lagged time series

```
1 df_lags = lagged_timeseries(stock_log, "CBA", 40)
2 df_lags
```

	T-40	T-39	T-38	T-37	T-36	T-35	T-34	T-33	T-32	T-31
Date										
1999-01-04	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1999-01-05	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
...	...	...	...	...	...	...	...	...	...	...
2025-06-26	0.021968	0.003894	0.014307	-0.016222	-0.001139	-0.004689	-0.002715	0.009202	0.000838	-0.006240
2025-06-27	0.003894	0.014307	-0.016222	-0.001139	-0.004689	-0.002715	0.009202	0.000838	-0.006240	0.008153

6744 rows × 41 columns

Split into training and testing

```

1 # Split the data in time
2 X_train = df_lags.loc[:"2021"]
3 X_val = df_lags.loc["2022-01":"2024-12"] # 2022-2024
4 X_test = df_lags.loc["2025":] # 2025
5
6 # Remove any with NAs and split into X and y
7 X_train = X_train.dropna()
8 X_val = X_val.dropna()
9 X_test = X_test.dropna()
10
11 y_train = X_train.pop("T")
12 y_val = X_val.pop("T")
13 y_test = X_test.pop("T")

```

```
1 X_train.shape, y_train.shape, X_val.shape, y_val.shape, X_test.shape, y_test.shape
```

```
((5825, 40), (5825,), (757, 40), (757,), (122, 40), (122,))
```


Inspect the split data

1 X_train

	T-40	T-39	T-38	T-37	T-36	T-35	T-34	T-33	T-32	T-31
Date										
1999-03-01	-0.010858	0.000437	0.008042	0.025859	-0.021323	-0.010573	-0.011214	-0.014823	-0.009704	0.018119
1999-03-02	0.000437	0.008042	0.025859	-0.021323	-0.010573	-0.011214	-0.014823	-0.009704	0.018119	0.012994
1999-03-03	0.008042	0.025859	-0.021323	-0.010573	-0.011214	-0.014823	-0.009704	0.018119	0.012994	0.010881
...	...	...	...	...	...	...	...	...	...	...
2021-12-29	0.015263	-0.004999	0.011656	0.013921	0.011090	0.003821	-0.012058	0.000919	-0.016106	0.008825
2021-12-30	-0.004999	0.011656	0.013921	0.011090	0.003821	-0.012058	0.000919	-0.016106	0.008825	-0.001018
2021-12-31	0.011656	0.013921	0.011090	0.003821	-0.012058	0.000919	-0.016106	0.008825	-0.001018	-0.003060

5825 rows × 40 columns

Plot the split




Fit a linear model

```
1 lr = LinearRegression()  
2 lr.fit(X_train, y_train);
```

Make a forecast for the test data:


```
1 y_pred = lr.predict(X_test)
```

Plot predictions in log return space




Convert to raw prices and plot

```
1 linear_log_forecasts = pd.Series(y_pred, index=y_test.index)
2 prev_price = stock.shift(1).loc["2025", "CBA"]
3 linear_price_forecasts = prev_price * np.exp(linear_log_forecasts)
```



Distribution of log return predictions



```
1 # Calculate MSE on raw prices using explicit forecast calculation
2 test_actual_prices = stock.loc["2025":, "CBA"]
3 test_linear_prices = linear_price_forecasts.loc["2025":]
4 linear_mse = mean_squared_error(test_actual_prices, test_linear_prices)
5 linear_mse
```

5.136664038344955

```
1 persistence_mse, trend_mse, linear_mse
```

(254.04075100411256, 99.28717915684173, 5.136664038344955)

Lecture Outline


- Time Series
- Baseline forecasts
- **Multi-step forecasts**
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series

Comparing apples to apples

The linear model is only producing *one-step-ahead* forecasts.

The other models are producing *multi-step-ahead* forecasts.

```
1 shifted_forecast = stock["CBA"].shift(1).loc["2025":]
```



```
1 shifted_mse = mean_squared_error(stock.loc["2025":, "CBA"], shifted_forecast)
2 persistence_mse, trend_mse, linear_mse, shifted_mse
```

(254.04075100411256, 99.28717915684173, 5.136664038344955, 5.06190042634739)

Autoregressive forecasts

The linear model needs the last 90 days to make a forecast.

Idea: Make the first forecast, then use that to make the next forecast, and so on.

$$\hat{y}_t = \beta_0 + \beta_1 y_{t-1} + \beta_2 y_{t-2} + \dots + \beta_n y_{t-n}$$

$$\hat{y}_{t+1} = \beta_0 + \beta_1 \hat{y}_t + \beta_2 y_{t-1} + \dots + \beta_n y_{t-n+1}$$

$$\hat{y}_{t+2} = \beta_0 + \beta_1 \hat{y}_{t+1} + \beta_2 \hat{y}_t + \dots + \beta_n y_{t-n+2}$$

$$\vdots$$

$$\hat{y}_{t+k} = \beta_0 + \beta_1 \hat{y}_{t+k-1} + \beta_2 \hat{y}_{t+k-2} + \dots + \beta_n \hat{y}_{t+k-n}$$

Autoregressive forecasting function


```

1  def autoregressive_forecast(model, X_val, suppress=False):
2      """
3      Generate a multi-step forecast using the given model.
4      """
5      multi_step = pd.Series(index=X_val.index, name="Multi Step")
6
7      # Initialize the input data for forecasting
8      input_data = X_val.iloc[0].values.reshape(1, -1)
9
10     for i in range(len(multi_step)):
11         # Ensure input_data has the correct feature names
12         input_df = pd.DataFrame(input_data, columns=X_val.columns)
13         if suppress:
14             next_value = model.predict(input_df, verbose=0)
15         else:
16             next_value = model.predict(input_df)
17
18         multi_step.iloc[i] = next_value.flatten()[0]
19
20         # Append that prediction to the input for the next forecast
21         if i + 1 < len(multi_step):
22             input_data = np.append(input_data[:, 1:], next_value).reshape(1, -1)
23
24     return multi_step


```

Look at the autoregressive linear forecasts

```
1 lr_forecast = autoregressive_forecast(lr, X_test)
2 lr_multi_step_prices = log_to_price(lr_forecast, last_price)
```



Compare log return forecasts



Metrics

One-step-ahead forecasts:

```
1 linear_mse, shifted_mse
```


```
(5.136664038344955, 5.06190042634739)
```

Multi-step-ahead forecasts:

```
1 multi_step_linear_mse = mean_squared_error(stock.loc["2025":, "CBA"], lr_multi_step_price  
2 persistence_mse, trend_mse, multi_step_linear_mse
```

```
(254.04075100411256, 99.28717915684173, 181.11397713761005)
```

Prefer only short windows



“It’s tough to make predictions, especially about the future.”

Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- **Neural network forecasts**
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series

Simple feedforward neural network

```

1 model = Sequential([
2     Rescaling(1/0.02),
3     Dense(32, activation="leaky_relu"),
4     Dense(1)]) # Linear activation for log returns
5 model.compile(optimizer="adam", loss="mean_absolute_error")

```

```

1 if Path("models/aus_fin_fnn_model.keras").exists():
2     model = keras.models.load_model("models/aus_fin_fnn_model.keras")
3 else:
4     es = EarlyStopping(patience=15, restore_best_weights=True)
5     model.fit(X_train, y_train, validation_data=(X_val, y_val), epochs=500,
6             callbacks=[es], verbose=0)
7     model.save("models/aus_fin_fnn_model.keras")
8
9 model.summary()

```

Model: "sequential"


Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(32, 40)	0
dense (Dense)	(32, 32)	1,312
dense_1 (Dense)	(32, 1)	33

Total params: 4,037 (15.77 KB)

Trainable params: 1,345 (5.25 KB)


Forecast and plot

```
1 y_pred = model.predict(X_test, verbose=0)
```




Plot forecasts in raw price space

```
1 fnn_log_forecasts = pd.Series(y_pred.flatten(), index=y_test.index)
2 fnn_price_forecasts = prev_price * np.exp(fnn_log_forecasts)
```




Distribution of log return predictions




Autoregressive forecasts

```
1 fnn_forecast = autoregressive_forecast(model, X_test, True)
2 fnn_multi_step_prices = log_to_price(fnn_forecast, last_price)
```



Compare multi-step log return forecasts



Metrics

One-step-ahead forecasts (MSE on raw prices):

```
1 nn_mse = mean_squared_error(stock.loc["2025":, "CBA"], fnn_price_forecasts.loc["2025":])  
2 linear_mse, nn_mse
```

(5.136664038344955, 5.249813642619517)

Multi-step-ahead forecasts (MSE on raw prices):

```
1 multi_step_fnn_mse = mean_squared_error(stock.loc["2025":, "CBA"], fnn_multi_step_prices.  
2 persistence_mse, trend_mse, multi_step_linear_mse, multi_step_fnn_mse
```

(254.04075100411256, 99.28717915684173, 181.11397713761005, 213.404767512833)

Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- **Recurrent Neural Networks**
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series

Basic facts of RNNs

- A recurrent neural network is a type of neural network that is designed to process sequences of data (e.g. time series, sentences).
- A recurrent neural network is any network that contains a recurrent layer.
- A recurrent layer is a layer that processes data in a sequence.
- An RNN can have one or more recurrent layers.
- Weights are shared over time; this allows the model to be used on arbitrary-length sequences.

Applications

- *Forecasting*: revenue forecast, weather forecast, predict disease rate from medical history, etc.
- *Classification*: given a time series of the activities of a visitor on a website, classify whether the visitor is a bot or a human.
- *Event detection*: given a continuous data stream, identify the occurrence of a specific event. Example: Detect utterances like “Hey Alexa” from an audio stream.
- *Anomaly detection*: given a continuous data stream, detect anything unusual happening. Example: Detect unusual activity on the corporate network.

Origin of the name of RNNs


A recurrence relation is an equation that expresses each element of a sequence as a function of the preceding ones. More precisely, in the case where only the immediately preceding element is involved, a recurrence relation has the form

$$u_n = \psi(n, u_{n-1}) \quad \text{for } n > 0.$$

Example: Factorial $n! = n(n-1)!$ for $n > 0$ given $0! = 1$.

Diagram of an RNN cell

The RNN processes each data in the sequence one by one, while keeping memory of what came before.



Schematic of a recurrent neural network. E.g. SimpleRNN, LSTM, or GRU.

A SimpleRNN cell





Diagram of a SimpleRNN cell.

All the outputs before the final one are often discarded.

LSTM internals



Neural Network
Layer



Pointwise
Operation



Vector
Transfer



Concatenate



Copy

Source: Christopher Olah (2015), [Understanding LSTM Networks](#), Colah's Blog.

GRU internals




Diagram of a GRU cell.

$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- **Stock prediction with recurrent networks**
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series

SimpleRNN

```

1 from keras.layers import SimpleRNN, Reshape
2 model = Sequential([
3     Rescaling(1/0.02),
4     Reshape((-1, 1)),
5     SimpleRNN(64, activation="tanh"),
6     Dense(1)]) # Linear activation for log returns
7 model.compile(optimizer="adam", loss="mean_absolute_error")

```

```

1 es = EarlyStopping(patience=15, restore_best_weights=True)
2 model.fit(X_train, y_train, validation_data=(X_val, y_val),
3     epochs=500, callbacks=[es], verbose=0)
4 model.summary()

```


Model: "sequential_1"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(32, 40)	0
reshape (Reshape)	(32, 40, 1)	0
simple_rnn (SimpleRNN)	(32, 64)	4,224
dense_2 (Dense)	(32, 1)	65

Total params: 12,869 (50.27 KB)
 Trainable params: 4,289 (16.75 KB)
 Non-trainable params: 0 (0.00 B)


Forecast and plot

```
1 y_pred = model.predict(X_test.to_numpy(), verbose=0)
```




Plot forecasts in raw price space

```
1 rnn_log_forecasts = pd.Series(y_pred.flatten(), index=y_test.index)
2 rnn_price_forecasts = prev_price * np.exp(rnn_log_forecasts)
```



Multi-step forecasts

```
1 rnn_forecast = autoregressive_forecast(model, X_test, True)
2 rnn_multi_step_prices = log_to_price(rnn_forecast, last_price)
```



Metrics

One-step-ahead forecasts (MSE on raw prices):

```
1 rnn_mse = mean_squared_error(stock.loc["2025":, "CBA"], rnn_price_forecasts.loc["2025":])
2 linear_mse, nn_mse, rnn_mse
```

(5.136664038344955, 5.249813642619517, 5.279663252855173)

Multi-step-ahead forecasts (MSE on raw prices):

```
1 multi_step_rnn_mse = mean_squared_error(stock.loc["2025":, "CBA"], rnn_multi_step_prices.
2 persistence_mse, trend_mse, multi_step_linear_mse, multi_step_fnn_mse, multi_step_rnn_mse
```

(254.04075100411256,
99.28717915684173,
181.11397713761005,
213.404767512833,
363.6276432551184)

GRU

```

1 from keras.layers import GRU
2
3 model = Sequential([
4     Rescaling(1/0.02),
5     Reshape((-1, 1)),
6     GRU(16, activation="tanh"),
7     Dense(1)]) # Linear activation for log returns
8 model.compile(optimizer="adam", loss="mean_absolute_error")

```

```

1 es = EarlyStopping(patience=15, restore_best_weights=True)
2 model.fit(X_train, y_train, validation_data=(X_val, y_val),
3     epochs=500, callbacks=[es], verbose=0)
4 model.summary()

```

Model: "sequential_2"


Layer (type)	Output Shape	Param #
rescaling_2 (Rescaling)	(32, 40)	0
reshape_1 (Reshape)	(32, 40, 1)	0
gru (GRU)	(32, 16)	912
dense_3 (Dense)	(32, 1)	17

Total params: 2,789 (10.89 KB)

Trainable params: 929 (3.63 KB)


Forecast and plot

```
1 y_pred = model.predict(X_test, verbose=0)
```




Plot forecasts in raw price space

```
1 gru_log_forecasts = pd.Series(y_pred.flatten(), index=y_test.index)
2 gru_price_forecasts = prev_price * np.exp(gru_log_forecasts)
```




Multi-step forecasts

```
1 gru_forecast = autoregressive_forecast(model, X_test, True)
2 gru_multi_step_prices = log_to_price(gru_forecast, last_price)
```



All multi-step log return forecasts



Metrics

One-step-ahead forecasts (MSE on raw prices):

```
1 gru_mse = mean_squared_error(stock.loc["2025":, "CBA"], gru_price_forecasts.loc["2025":])
2 linear_mse, nn_mse, rnn_mse, gru_mse
```

(5.136664038344955, 5.249813642619517, 5.279663252855173, 5.167950295726733)

Multi-step-ahead forecasts (MSE on raw prices):

```
1 multi_step_gru_mse = mean_squared_error(stock.loc["2025":, "CBA"], gru_multi_step_prices.
2 persistence_mse, trend_mse, multi_step_linear_mse, multi_step_fnn_mse, multi_step_rnn_mse
```

(254.04075100411256,
99.28717915684173,
181.11397713761005,
213.404767512833,
363.6276432551184,
94.02476396370041)

Summary of all model performance

One-step forecasts

	MSE
Shifted	5.061900
Linear	5.136664
FNN	5.249814
SimpleRNN	5.279663
GRU	5.167950

Multi-step forecasts

	MSE
Persistence	254.040751
Trend	99.287179
Linear	181.113977
FNN	213.404768
SimpleRNN	363.627643
GRU	94.024764

Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- **Internals of the SimpleRNN**
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series

The rank of a time series

Say we had n observations of a time series $x_1, x_2, \dots, x_n$.

This $\mathbf{x} = (x_1, \dots, x_n)$ would have shape $(n,)$ & rank 1.

If instead we had a batch of b time series'

$$\mathbf{X} = \begin{pmatrix} x_7 & x_8 & \dots & x_{7+n-1} \\ x_2 & x_3 & \dots & x_{2+n-1} \\ \vdots & \vdots & \ddots & \vdots \\ x_3 & x_4 & \dots & x_{3+n-1} \end{pmatrix},$$

the batch $\mathbf{X}$ would have shape (b, n) & rank 2.

Multivariate time series

t	x	y
0	x_0	y_0
1	x_1	y_1
2	x_2	y_2
3	x_3	y_3

Say n observations of the m time series, would be a shape (n, m) matrix of rank 2.

In Keras, a batch of b of these time series has shape (b, n, m) and has rank 3.

i Note

Use $\mathbf{x}_t \in \mathbb{R}^{1 \times m}$ to denote the vector of all time series at time t . Here, $\mathbf{x}_t = (x_t, y_t)$.

SimpleRNN

Say each prediction is a vector of size d , so $\mathbf{y}_t \in \mathbb{R}^{1 \times d}$.

Then the main equation of a SimpleRNN, given $\mathbf{y}_0 = \mathbf{0}$, is

$$\mathbf{y}_t = \psi(\mathbf{x}_t \mathbf{W}_x + \mathbf{y}_{t-1} \mathbf{W}_y + \mathbf{b}).$$

Here,

$$\begin{aligned} \mathbf{x}_t &\in \mathbb{R}^{1 \times m}, \mathbf{W}_x \in \mathbb{R}^{m \times d}, \\ \mathbf{y}_{t-1} &\in \mathbb{R}^{1 \times d}, \mathbf{W}_y \in \mathbb{R}^{d \times d}, \text{ and } \mathbf{b} \in \mathbb{R}^d. \end{aligned}$$

SimpleRNN (in batches)

Say we operate on batches of size b , then $\mathbf{Y}_t \in \mathbb{R}^{b \times d}$.

The main equation of a SimpleRNN, given $\mathbf{Y}_0 = \mathbf{0}$, is

$$\mathbf{Y}_t = \psi(\mathbf{X}_t \mathbf{W}_x + \mathbf{Y}_{t-1} \mathbf{W}_y + \mathbf{b}).$$

Here,

$$\begin{aligned} \mathbf{X}_t &\in \mathbb{R}^{b \times m}, \mathbf{W}_x \in \mathbb{R}^{m \times d}, \\ \mathbf{Y}_{t-1} &\in \mathbb{R}^{b \times d}, \mathbf{W}_y \in \mathbb{R}^{d \times d}, \text{ and } \mathbf{b} \in \mathbb{R}^d. \end{aligned}$$

Remember, $\mathbf{X} \in \mathbb{R}^{b \times n \times m}$, $\mathbf{Y} \in \mathbb{R}^{b \times d}$, and $\mathbf{X}_t$ is equivalent to $\mathbf{X}[:, t, :]$.



Simple Keras demo

```

1 num_obs = 4
2 num_time_steps = 3
3 num_time_series = 2
4
5 X = (
6     np.arange(num_obs * num_time_steps * num_time_series)
7     .astype(np.float32)
8     .reshape([num_obs, num_time_steps, num_time_series])
9 )
10
11 output_size = 1
12 y = np.array([0, 0, 1, 1])

```

```
1 X[:2]
```

```

array([[[ 0.,  1.],
        [ 2.,  3.],
        [ 4.,  5.]],

       [[ 6.,  7.],
        [ 8.,  9.],
        [10., 11.]]], dtype=float32)

```

```
1 X[2:]
```

```

array([[[12., 13.],
        [14., 15.],
        [16., 17.]],

       [[18., 19.],
        [20., 21.],
        [22., 23.]]], dtype=float32)

```


Keras' SimpleRNN

As usual, the `SimpleRNN` is just a layer in Keras.

```
1 from keras.layers import SimpleRNN
2
3 random.seed(1234)
4 model = Sequential([SimpleRNN(output_size, activation="sigmoid")])
5 model.compile(loss="binary_crossentropy", metrics=["accuracy"])
6
7 hist = model.fit(X, y, epochs=500, verbose=False)
8 model.evaluate(X, y, verbose=False)
```

```
[8.05906867980957, 0.5]
```

The predicted probabilities on the training set are:

```
1 model.predict(X, verbose=0)
```

```
array([[8.56e-05],
       [2.25e-10],
       [5.98e-16],
       [1.59e-21]], dtype=float32)
```

SimpleRNN weights

```
1 model.get_weights()
```

```
[array([[ -1.47],
        [-0.67]], dtype=float32),
 array([[0.99]], dtype=float32),
 array([-0.14], dtype=float32)]
```


```
1 def sigmoid(x):
2     return 1 / (1 + np.exp(-x))
3
4
5 W_x, W_y, b = model.get_weights()
6
7 Y = np.zeros((num_obs, output_size), dtype=np.float32)
8 for t in range(num_time_steps):
9     X_t = X[:, t, :]
10    z = X_t @ W_x + Y @ W_y + b
11    Y = sigmoid(z)
12
13 Y
```

```
array([[8.56e-05],
        [2.25e-10],
        [5.98e-16],
        [1.59e-21]], dtype=float32)
```


Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- **Other recurrent network variants**
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- Predicting Multiple Time Series


Input and output sequences




(a) Many to many



(b) Many to one



(c) One to Many



(d) Encoder-decoder

Categories of recurrent neural networks: sequence to sequence, sequence to vector, vector to sequence, encoder-decoder network.

Source: Melissa Renard (2025)


Input and output sequences

- Sequence to sequence: Useful for predicting time series such as using prices over the last N days to output the prices shifted one day into the future (i.e. from $N - 1$ days ago to tomorrow.)
- Sequence to vector: ignore all outputs in the previous time steps except for the last one. Example: give a sentiment score to a sequence of words corresponding to a movie review.

Input and output sequences

- Vector to sequence: feed the network the same input vector over and over at each time step and let it output a sequence. Example: given that the input is an image, find a caption for it. The image is treated as an input vector (pixels in an image do not follow a sequence). The caption is a sequence of textual description of the image. A dataset containing images and their descriptions is the input of the RNN.
- The Encoder-Decoder: The encoder is a sequence-to-vector network. The decoder is a vector-to-sequence network. Example: Feed the network a sequence in one language. Use the encoder to convert the sentence into a single vector representation. The decoder decodes this vector into the translation of the sentence in another language.

Recurrent layers can be stacked.



Deep RNN unrolled through time.

Source: Melissa Renard (2025)

Beyond RNNs

RNNs are **foundational** and the concepts transfer everywhere, but in practice they have been largely overtaken by newer architectures:

- **Attention / Transformers** (Vaswani et al., 2017): instead of processing a sequence step-by-step, attention lets every position look at every other position simultaneously. This avoids the vanishing gradient problem over long sequences and enables massive parallelism during training.
- **State-space models** (e.g. Mamba, 2023): process sequences in linear time rather than the quadratic cost of full attention, while still capturing long-range dependencies.


Why learn RNNs then?

- The **hidden state** concept carries directly into transformers (which use it differently) and state-space models
- **Gating** (LSTM/GRU) was a key insight that influenced later architectures
- RNNs remain a strong, simple baseline for short sequences and small datasets
- Understanding RNNs makes it easier to appreciate *why* attention was such a breakthrough

Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- **CoreLogic Hedonic Home Value Index**
- Predicting Sydney House Prices
- Predicting Multiple Time Series

Australian House Price Indices



i Note

I apologise in advance for not being able to share this dataset with anyone (it is not mine to share).

Percentage changes


```
1 changes = house_prices.pct_change().dropna()
2 changes.round(2)
```

	Brisbane	East_Bris	North_Bris	West_Bris	Melbourne	North_Syd	Sydney
Date							
1990-02-28	0.03	-0.01	0.01	0.01	0.00	-0.00	-0.02
1990-03-31	0.01	0.03	0.01	0.01	0.02	-0.00	0.03
1990-04-30	0.02	0.02	0.01	-0.00	0.01	0.03	0.04
...	...	...	...	...	...	...	...
2021-03-31	0.04	0.04	0.03	0.04	0.02	0.05	0.05
2021-04-30	0.03	0.01	0.01	-0.00	0.01	0.02	0.02
2021-05-31	0.03	0.03	0.03	0.03	0.03	0.02	0.04

376 rows × 7 columns

Percentage changes

```
1 changes.plot();
```



The size of the changes

```
1 changes.mean()
```


```
Brisbane      0.005496
East_Bris     0.005416
North_Bris    0.005024
West_Bris     0.004842
Melbourne     0.005677
North_Syd     0.004819
Sydney        0.005526
dtype: float64
```

```
1 changes *= 100
```

```
1 changes.mean()
```

```
Brisbane      0.549605
East_Bris     0.541562
North_Bris    0.502390
West_Bris     0.484204
Melbourne     0.567700
North_Syd     0.481863
Sydney        0.552641
dtype: float64
```


```
1 changes.plot(legend=False);
```



Split *without* shuffling

```
1 num_train = int(0.6 * len(changes))  
2 num_val = int(0.2 * len(changes))  
3 num_test = len(changes) - num_train - num_val  
4 print(f"# Train: {num_train}, # Val: {num_val}, # Test: {num_test}")
```

Train: 225, # Val: 75, # Test: 76



Subsequences of a time series

We can use numpy to convert a time series into subsequences/chunks using a sliding window.

```
1 integers = np.arange(10)
2 seq_len = 3
3
4 X = np.stack([integers[i : i + seq_len] for i in range(len(integers) - seq_len)])
5 y = integers[seq_len:]
6
7 for i in range(len(X)):
8     print(X[i].tolist(), int(y[i]))
```

```
[0, 1, 2] 3
[1, 2, 3] 4
[2, 3, 4] 5
[3, 4, 5] 6
[4, 5, 6] 7
[5, 6, 7] 8
[6, 7, 8] 9
```


Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- **Predicting Sydney House Prices**
- Predicting Multiple Time Series

Creating input arrays

```

1 def make_sequences(data, targets, seq_len, start_index=0, end_index=None):
2     data = np.asarray(data, dtype=np.float32)
3     targets = np.asarray(targets, dtype=np.float32)
4     if end_index is None:
5         end_index = len(data)
6     n = end_index - start_index - seq_len + 1
7     X = np.stack([data[start_index + i : start_index + i + seq_len] for i in range(n)])
8     y = targets[start_index : start_index + n]
9     return X, y

```

```

1 # Num. of input time series.
2 num_ts = changes.shape[1]
3
4 # How many prev. months to use.
5 seq_length = 6
6
7 # Predict the next month ahead.
8 ahead = 1
9
10 # The index of the first target.
11 delay = seq_length + ahead - 1

```

```

1 # Which suburb to predict.
2 target_suburb = changes["Sydney"]
3
4 X_train, y_train = make_sequences(
5     changes[:-delay], target_suburb[delay:
6     seq_length, end_index=num_train,
7 )

```

```

1 X_val, y_val = make_sequences(
2     changes[:-delay], target_suburb[delay:
3     seq_length, start_index=num_train,
4     end_index=num_train + num_val].

```

```

1 X_test, y_test = make_sequences(
2     changes[:-delay], target_suburb[delay:
3     seq_length, start_index=num_train + r
4 )

```

Training set shape

```
1 X_train.shape
```

```
(220, 6, 7)
```

```
1 y_train.shape
```

```
(220,)
```

A dense network

```

1 from keras.layers import Input, Flatten
2 random.seed(1)
3 model_dense = Sequential([
4     Input((seq_length, num_ts)),
5     Flatten(),
6     Dense(50, activation="leaky_relu"),
7     Dense(20, activation="leaky_relu"),
8     Dense(1, activation="linear")
9 ])
10 model_dense.compile(loss="mse", optimizer="adam")
11 print(f"This model has {model_dense.count_params()} parameters.")
12
13 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
14 %time hist = model_dense.fit(X_train, y_train, epochs=1_000, \
15     validation_data=(X_val, y_val), callbacks=[es], verbose=0);

```

This model has 3191 parameters.

Epoch 76: early stopping


Restoring model weights from the end of the best epoch: 26.

CPU times: user 1.36 s, sys: 81.9 ms, total: 1.44 s

Wall time: 1.38 s

Plot the model


```
1 from keras.utils import plot_model
2
3 plot_model(model_dense, show_shapes=True)
```



Assess the fits

```
1 model_dense.evaluate(X_val, y_val, verbose=0)
```

1.3529319763183594



A SimpleRNN layer

```

1 random.seed(1)
2
3 model_simple = Sequential([
4     Input((seq_length, num_ts)),
5     SimpleRNN(50),
6     Dense(1, activation="linear")
7 ])
8 model_simple.compile(loss="mse", optimizer="adam")
9 print(f"This model has {model_simple.count_params()} parameters.")
10
11 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
12 %time hist = model_simple.fit(X_train, y_train, epochs=1_000, \
13     validation_data=(X_val, y_val), callbacks=[es], verbose=0);

```

This model has 2951 parameters.

Epoch 53: early stopping


Restoring model weights from the end of the best epoch: 3.

CPU times: user 1.7 s, sys: 46.8 ms, total: 1.75 s

Wall time: 1.71 s

Plot the model


```
1 plot_model(model_simple, show_shapes=True)
```



Assess the fits

```
1 model_simple.evaluate(X_val, y_val, verbose=0)
```

0.8649981617927551



A LSTM layer

```

1 from keras.layers import LSTM
2
3 random.seed(1)
4
5 model_lstm = Sequential([
6     Input((seq_length, num_ts)),
7     LSTM(50),
8     Dense(1, activation="linear")
9 ])
10
11 model_lstm.compile(loss="mse", optimizer="adam")
12
13 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
14
15 %time hist = model_lstm.fit(X_train, y_train, epochs=1_000, \
16     validation_data=(X_val, y_val), callbacks=[es], verbose=0);

```

Epoch 58: early stopping

Restoring model weights from the end of the best epoch: 8.


CPU times: user 1.43 s, sys: 400 ms, total: 1.83 s

Wall time: 1.55 s

Assess the fits

```
1 model_lstm.evaluate(X_val, y_val, verbose=0)
```

0.8229678273200989



A GRU layer

```

1 from keras.layers import GRU
2
3 random.seed(1)
4
5 model_gru = Sequential([
6     Input((seq_length, num_ts)),
7     GRU(50),
8     Dense(1, activation="linear")
9 ])
10
11 model_gru.compile(loss="mse", optimizer="adam")
12
13 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
14
15 %time hist = model_gru.fit(X_train, y_train, epochs=1_000, \
16     validation_data=(X_val, y_val), callbacks=[es], verbose=0)

```

Epoch 59: early stopping

Restoring model weights from the end of the best epoch: 9.


CPU times: user 2.49 s, sys: 48.1 ms, total: 2.54 s

Wall time: 2.51 s

Assess the fits

```
1 model_gru.evaluate(X_val, y_val, verbose=0)
```

0.813169002532959



Two GRU layers

```

1 random.seed(1)
2
3 model_two_grus = Sequential([
4     Input((seq_length, num_ts)),
5     GRU(50, return_sequences=True),
6     GRU(50),
7     Dense(1, activation="linear")
8 ])
9
10 model_two_grus.compile(loss="mse", optimizer="adam")
11
12 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
13
14 %time hist = model_two_grus.fit(X_train, y_train, epochs=1_000, \
15     validation_data=(X_val, y_val), callbacks=[es], verbose=0)

```

Epoch 57: early stopping

Restoring model weights from the end of the best epoch: 7.


CPU times: user 4.32 s, sys: 102 ms, total: 4.42 s

Wall time: 4.37 s

Assess the fits

```
1 model_two_grus.evaluate(X_val, y_val, verbose=0)
```

0.7832801342010498



Compare the models


	Model	MSE
0	Dense	1.352932
1	SimpleRNN	0.864998
2	LSTM	0.822968
3	GRU	0.813169
4	2 GRUs	0.783280

The network with two GRU layers is the best.

```
1 model_two_grus.evaluate(X_test, y_test, verbose=0)
```

```
1.9257299900054932
```


Test set



Lecture Outline

- Time Series
- Baseline forecasts
- Multi-step forecasts
- Neural network forecasts
- Recurrent Neural Networks
- Stock prediction with recurrent networks
- Internals of the SimpleRNN
- Other recurrent network variants
- CoreLogic Hedonic Home Value Index
- Predicting Sydney House Prices
- **Predicting Multiple Time Series**

Creating input arrays

Change the `targets` argument to include all the suburbs.

```
1 X_val, y_val = make_sequences(  
2     changes[:-delay], changes[delay:],  
3     seq_length, start_index=num_train,  
4     end_index=num_train + num_val,  
5 )
```

```
1 X_train, y_train = make_sequences(  
2     changes[:-delay], changes[delay:],  
3     seq_length, end_index=num_train,  
4 )
```

```
1 X_test, y_test = make_sequences(  
2     changes[:-delay], changes[delay:],  
3     seq_length, start_index=num_train + r  
4 )
```

Training set shape

The shape of our training set is now:

```
1 X_train.shape
```

```
(220, 6, 7)
```

```
1 y_train.shape
```

```
(220, 7)
```

A dense network

```

1 random.seed(1)
2 model_dense = Sequential([
3     Input((seq_length, num_ts)),
4     Flatten(),
5     Dense(50, activation="leaky_relu"),
6     Dense(20, activation="leaky_relu"),
7     Dense(num_ts, activation="linear")
8 ])
9 model_dense.compile(loss="mse", optimizer="adam")
10 print(f"This model has {model_dense.count_params()} parameters.")
11
12 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
13 %time hist = model_dense.fit(X_train, y_train, epochs=1_000, \
14     validation_data=(X_val, y_val), callbacks=[es], verbose=0);

```

This model has 3317 parameters.

Epoch 70: early stopping


Restoring model weights from the end of the best epoch: 20.

CPU times: user 1.14 s, sys: 82.6 ms, total: 1.22 s

Wall time: 1.17 s

Plot the model





```
1 plot_model(model_dense, show_shapes=True)
```



Assess the fits

```
1 model_dense.evaluate(X_val, y_val, verbose=0)
```

1.4135159254074097



A SimpleRNN layer

```

1 random.seed(1)
2
3 model_simple = Sequential([
4     Input((seq_length, num_ts)),
5     SimpleRNN(50),
6     Dense(num_ts, activation="linear")
7 ])
8 model_simple.compile(loss="mse", optimizer="adam")
9 print(f"This model has {model_simple.count_params()} parameters.")
10
11 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
12 %time hist = model_simple.fit(X_train, y_train, epochs=1_000, \
13     validation_data=(X_val, y_val), callbacks=[es], verbose=0);

```

This model has 3257 parameters.

Epoch 69: early stopping


Restoring model weights from the end of the best epoch: 19.

CPU times: user 2.2 s, sys: 68.3 ms, total: 2.27 s

Wall time: 2.22 s

Plot the model





```
1 plot_model(model_simple, show_shapes=True)
```



Assess the fits

```
1 model_simple.evaluate(X_val, y_val, verbose=0)
```

1.6102873086929321



A LSTM layer

```

1 random.seed(1)
2
3 model_lstm = Sequential([
4     Input((seq_length, num_ts)),
5     LSTM(50),
6     Dense(num_ts, activation="linear")
7 ])
8
9 model_lstm.compile(loss="mse", optimizer="adam")
10
11 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
12
13 %time hist = model_lstm.fit(X_train, y_train, epochs=1_000, \
14     validation_data=(X_val, y_val), callbacks=[es], verbose=0);

```

Epoch 76: early stopping

Restoring model weights from the end of the best epoch: 26.





CPU times: user 1.87 s, sys: 540 ms, total: 2.41 s

Wall time: 2.03 s

Assess the fits

```
1 model_lstm.evaluate(X_val, y_val, verbose=0)
```

1.3229057788848877



A GRU layer

```

1 random.seed(1)
2
3 model_gru = Sequential([
4     Input((seq_length, num_ts)),
5     GRU(50),
6     Dense(num_ts, activation="linear")
7 ])
8
9 model_gru.compile(loss="mse", optimizer="adam")
10
11 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
12
13 %time hist = model_gru.fit(X_train, y_train, epochs=1_000, \
14     validation_data=(X_val, y_val), callbacks=[es], verbose=0)

```

Epoch 69: early stopping

Restoring model weights from the end of the best epoch: 19.





CPU times: user 3.06 s, sys: 72.6 ms, total: 3.13 s

Wall time: 3.09 s

Assess the fits

```
1 model_gru.evaluate(X_val, y_val, verbose=0)
```

1.3403815031051636



Two GRU layers

```
1 random.seed(1)
2
3 model_two_grus = Sequential([
4     Input((seq_length, num_ts)),
5     GRU(50, return_sequences=True),
6     GRU(50),
7     Dense(num_ts, activation="linear")
8 ])
9
10 model_two_grus.compile(loss="mse", optimizer="adam")
11
12 es = EarlyStopping(patience=50, restore_best_weights=True, verbose=1)
13
14 %time hist = model_two_grus.fit(X_train, y_train, epochs=1_000, \
15     validation_data=(X_val, y_val), callbacks=[es], verbose=0)
```

Epoch 74: early stopping

Restoring model weights from the end of the best epoch: 24.





CPU times: user 5.4 s, sys: 123 ms, total: 5.52 s

Wall time: 5.44 s

Assess the fits

```
1 model_two_grus.evaluate(X_val, y_val, verbose=0)
```

1.3471126556396484



Compare the models




	Model	MSE
1	SimpleRNN	1.610287
0	Dense	1.413516
4	2 GRUs	1.347113
3	GRU	1.340382
2	LSTM	1.322906

The network with an LSTM layer is the best.

```
1 model_lstm.evaluate(X_test, y_test, verbose=0)
```

```
1.8732062578201294
```

Test set



Package Versions

```
1 from watermark import watermark
2 print(watermark(python=True, packages="keras,matplotlib,numpy,pandas,seaborn,scipy,torch")
```

Python implementation: CPython

Python version : 3.14.5

IPython version : 9.13.0

keras : 3.14.1

matplotlib: 3.10.9

numpy : 2.4.4

pandas : 3.0.2

seaborn : 0.13.2

scipy : 1.17.1

torch : 2.11.0

Glossary

- autoregressive forecasting
- forecasting
- GRU
- LSTM
- one-step/multi-step ahead forecasting
- persistence forecast
- recurrent neural networks
- SimpleRNN